

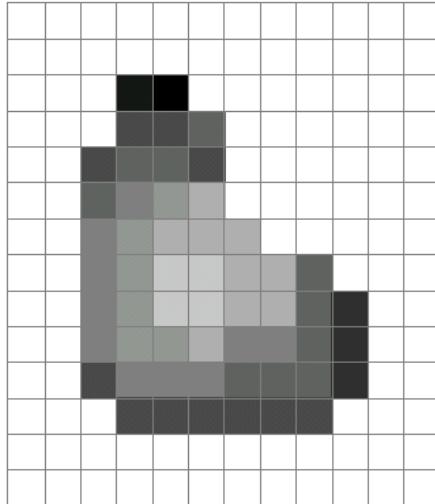
# CSE 6367: Computer Vision

## Image Filtering and Edge Detection

Slide Courtesy: Richard  
Szeliski, Computer Vision:  
Algorithms and Applications,  
2nd Edition

# What is an image?

- A grid (matrix) of intensity values

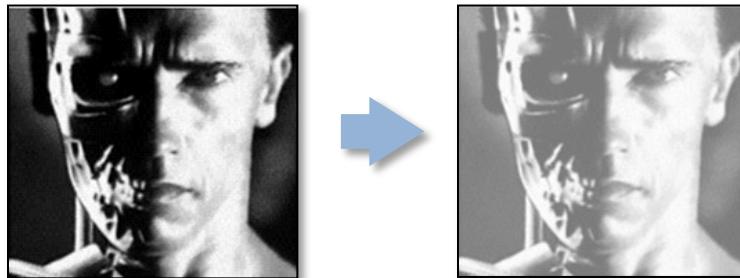


255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	20	0	255	255	255	255	255	255	255	255	255	255	255
255	255	255	75	75	75	255	255	255	255	255	255	255	255	255	255
255	255	75	95	95	75	255	255	255	255	255	255	255	255	255	255
255	255	96	127	145	175	255	255	255	255	255	255	255	255	255	255
255	255	127	145	175	175	175	255	255	255	255	255	255	255	255	255
255	255	127	145	200	200	175	175	95	255	255	255	255	255	255	255
255	255	127	145	200	200	175	175	95	47	255	255	255	255	255	255
255	255	127	145	145	175	127	127	95	47	255	255	255	255	255	255
255	255	74	127	127	127	95	95	95	47	255	255	255	255	255	255
255	255	255	74	74	74	74	74	74	74	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255
255	255	255	255	255	255	255	255	255	255	255	255	255	255	255	255

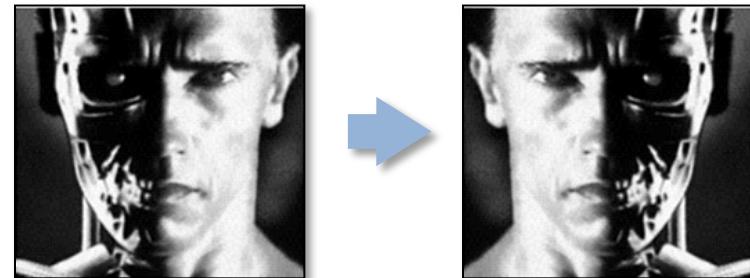
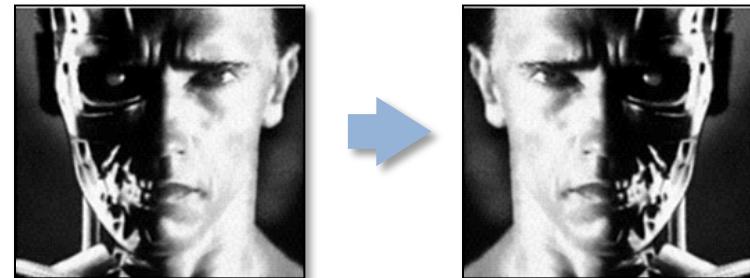
(common to use one byte per value: 0 = black, 255 = white)

# Image transformations

- As with any function, we can apply operators to an image



$$g(x,y) = f(x,y) + 20$$



$$g(x,y) = f(-x,y)$$

- Today we'll talk about a special kind of operator, *convolution* (linear filtering)

# Filters

- Filtering
  - Form a new image whose pixel values are a combination of the original pixel values
- Why?
  - To get useful information from images
    - E.g., extract edges or contours (to understand shape)
  - To enhance the image
    - E.g., to remove noise
    - E.g., to sharpen and “enhance image” a la CSI
  - A key operator in Convolutional Neural Networks

# Canonical Image Processing problems

- Image Restoration
  - denoising
  - deblurring
- Image Compression
  - JPEG, HEIF, MPEG, ...
- Locating Structural Features
  - corners
  - edges

# Question: Noise reduction

- Given a camera and a still scene, how can you reduce noise?



Take lots of images and average them!

What's the next best thing?

Source: S. Seitz

# Image filtering

- Modify the pixels in an image based on some function of a local neighborhood of each pixel

10	5	3
4	5	1
1	1	7

Local image data

Some function

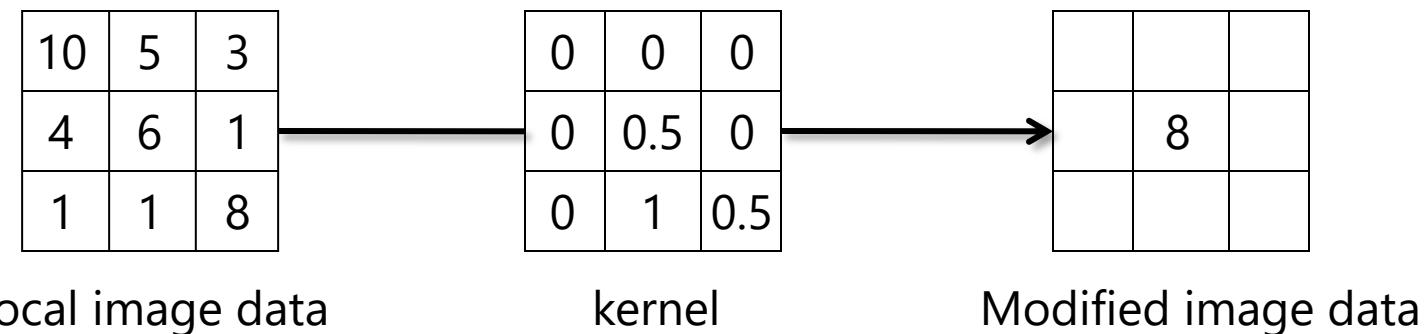


	7	

Modified image data

# Linear filtering

- One simple version of filtering: linear filtering (cross-correlation, convolution)
  - Replace each pixel by a linear combination (a weighted sum) of its neighbors
- The prescription for the linear combination is called the “kernel” (or “mask”, “filter”)



# Cross-correlation

Let  $F$  be the image,  $H$  be the kernel (of size  $2k+1 \times 2k+1$ ),  
and  $G$  be the output image

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

This is called a **cross-correlation** operation:

$$G = H \otimes F$$

- Can think of as a “dot product” between local neighborhood and kernel for each pixel

# Convolution

- Same as cross-correlation, except that the kernel is “flipped” (horizontally and vertically)

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i - u, j - v]$$

This is called a **convolution** operation:

$$G = H * F$$

- Convolution is **commutative** and **associative**

# Mean filtering




$H$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0



$F$

	0	10	20	30	30	30	20	10		
	0	20	40	60	60	60	40	20		
	0	30	60	90	90	90	60	30		
	0	30	50	80	80	90	60	30		
	0	30	50	80	80	90	60	30		
	0	20	30	50	50	60	40	20		
10	20	30	30	30	30	20	10			
10	10	10	0	0	0	0	0	0		

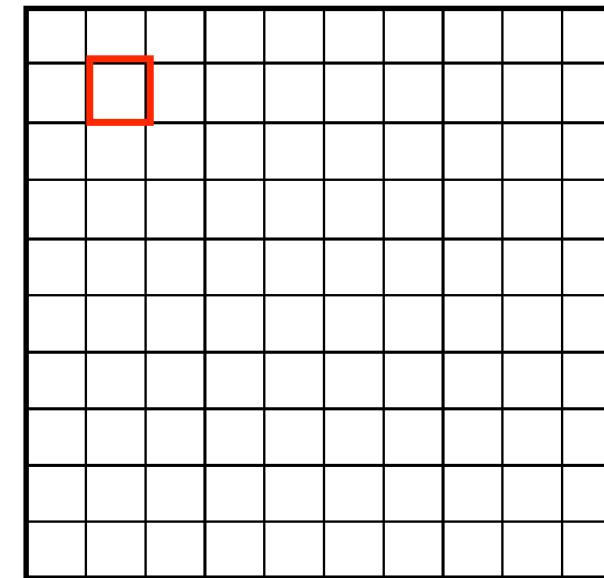
$G$

# Mean filtering/Moving average

$$F[x, y]$$

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	0
0	0	0	90	90	90	90	0
0	0	0	90	90	90	90	0
0	0	0	90	0	90	90	0
0	0	0	90	90	90	90	0
0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0
0	0	0	0	0	0	0	0

$$G[x, y]$$



# Mean filtering/Moving average

$$F[x, y]$$

$$G[x, y]$$

# Mean filtering/Moving average

$$F[x, y]$$

$$G[x, y]$$

# Mean filtering/Moving average

$$F[x, y]$$

$$G[x, y]$$

# Mean filtering/Moving average

$F[x, y]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$G[x, y]$

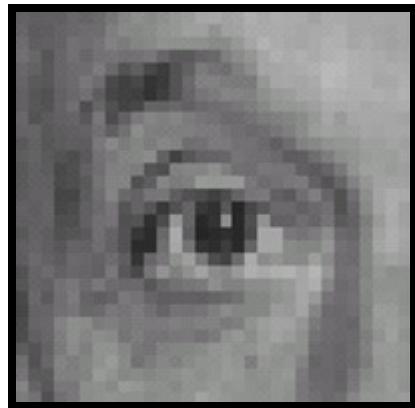

0    10    20    30    30

# Mean filtering/Moving average

$$F[x, y]$$

$$G[x, y]$$

# Linear filters: examples

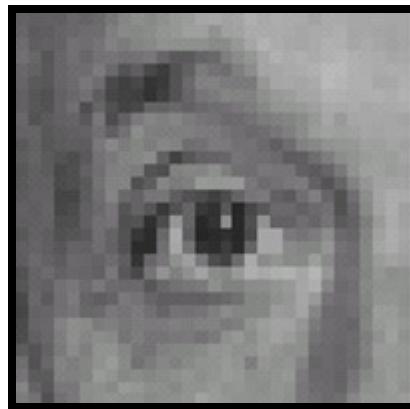


\*

0	0	0
0	1	0
0	0	0

Original

# Linear filters: examples

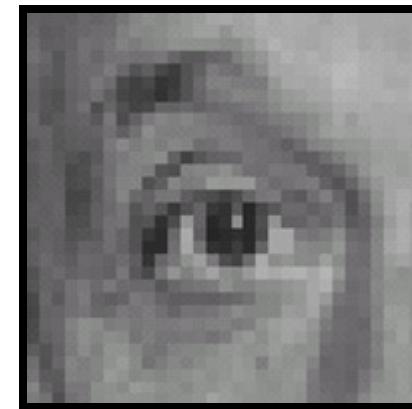


Original

\*

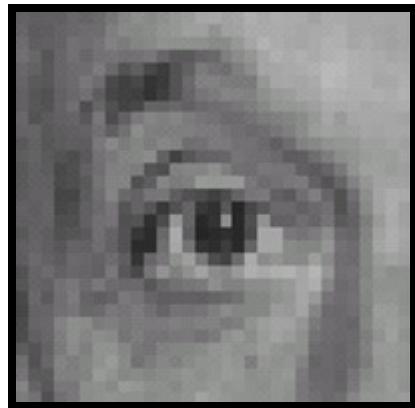
0	0	0
0	1	0
0	0	0

=



Identical image

# Linear filters: examples



\*

0	0	0
1	0	0
0	0	0

# Linear filters: examples

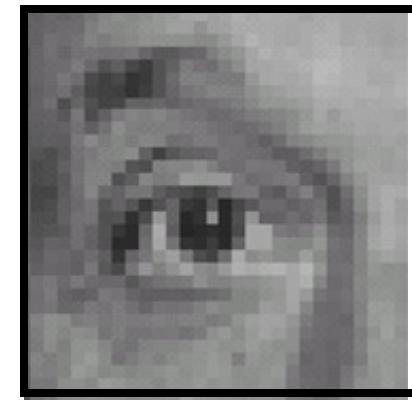


Original

\*

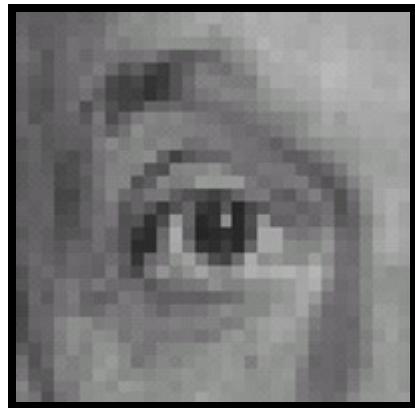
0	0	0
1	0	0
0	0	0

=



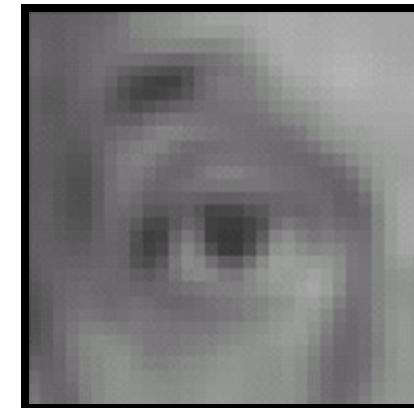
Shifted left by 1 pixel

# Linear filters: examples



Original

$$\ast \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} =$$



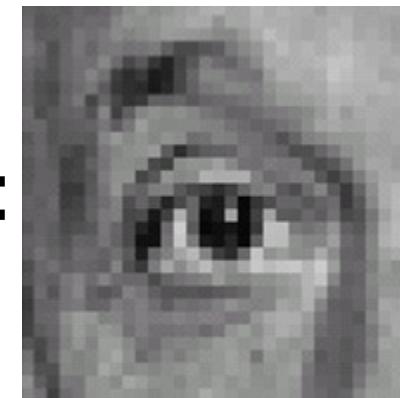
Blur (with a mean filter)

# Linear filters: examples



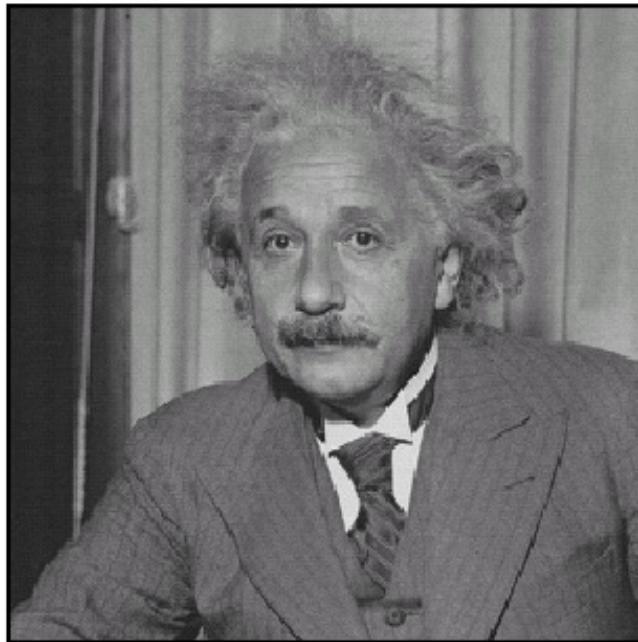
Original

$$\ast \left( \begin{array}{|ccc|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} - \frac{1}{9} \begin{array}{|ccc|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} \right) =$$

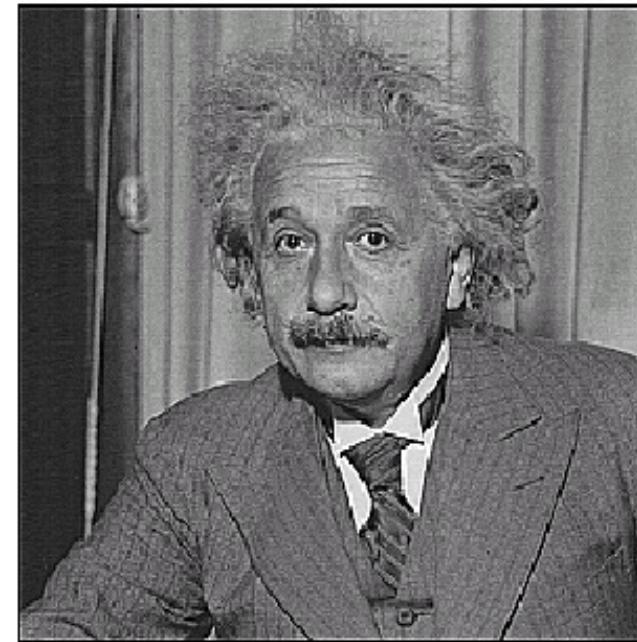


**Sharpening filter**  
(accentuates edges)

# Sharpening

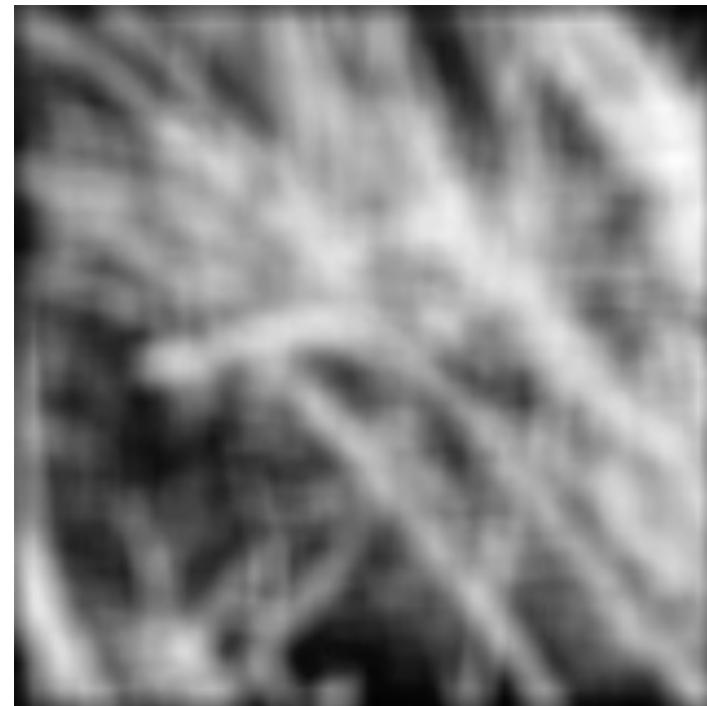
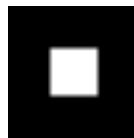


**before**

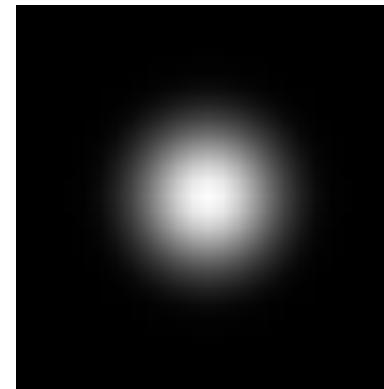
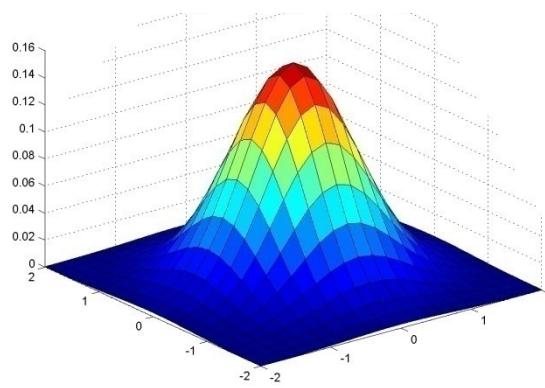


**after**

# Smoothing with box filter revisited

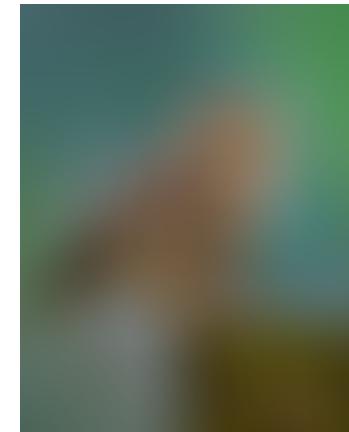
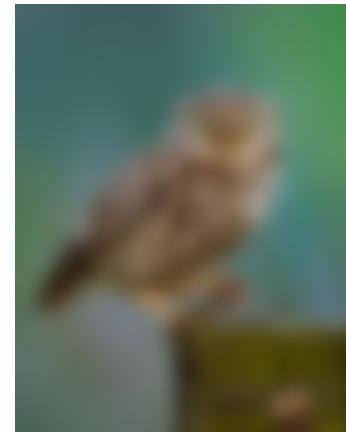


# Gaussian kernel

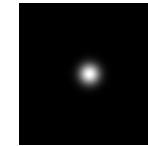


$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

# Gaussian filters



$\sigma = 1$  pixel



$\sigma = 5$  pixels

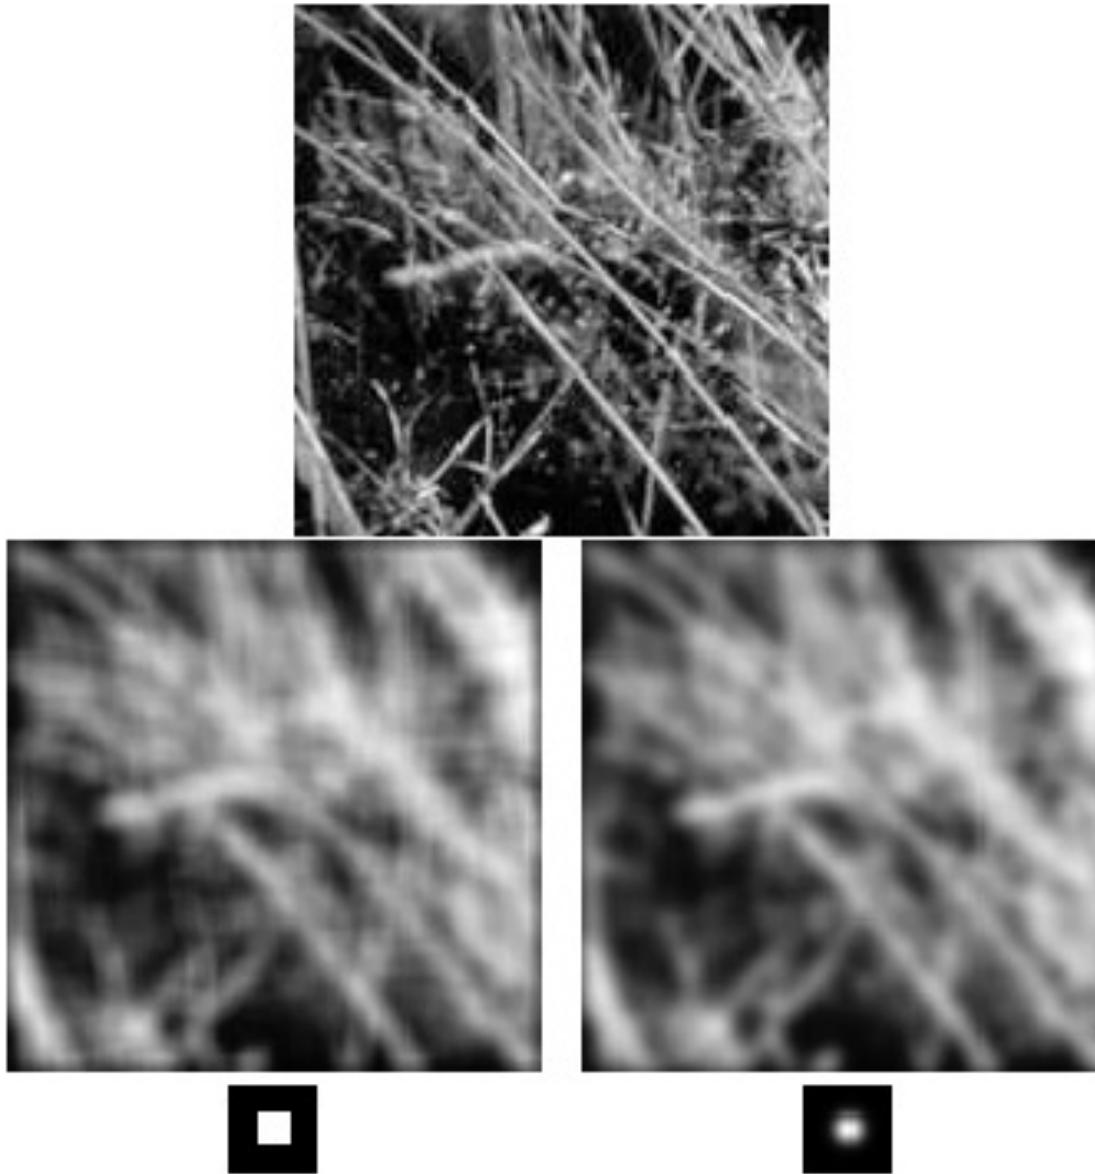


$\sigma = 10$  pixels



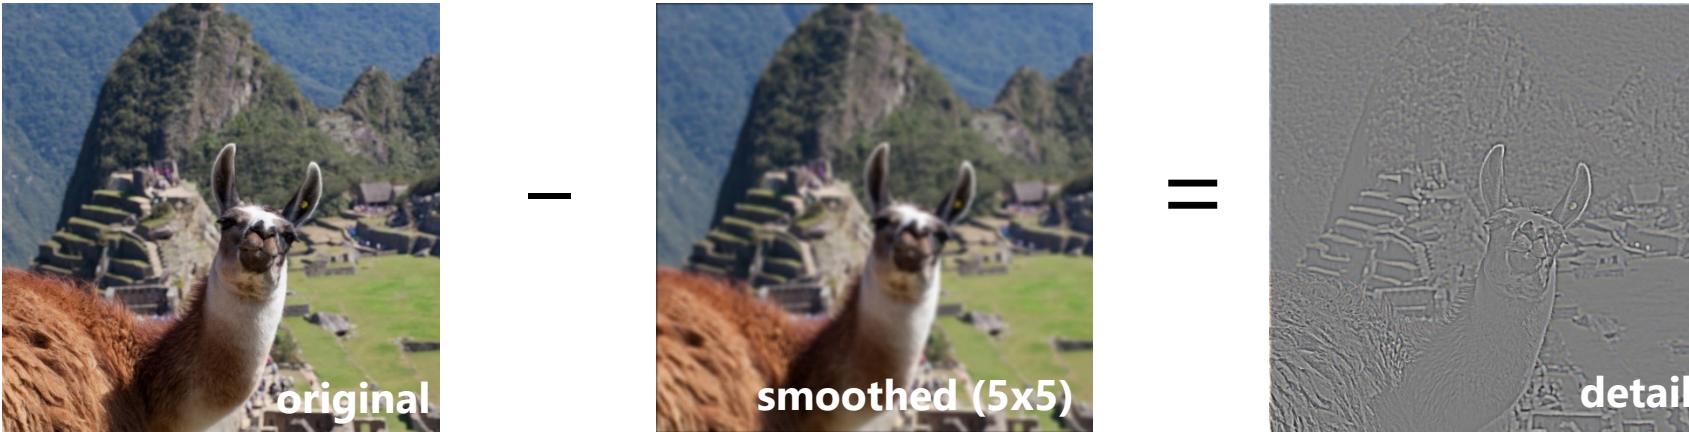
$\sigma = 30$  pixels

# Mean vs. Gaussian filtering



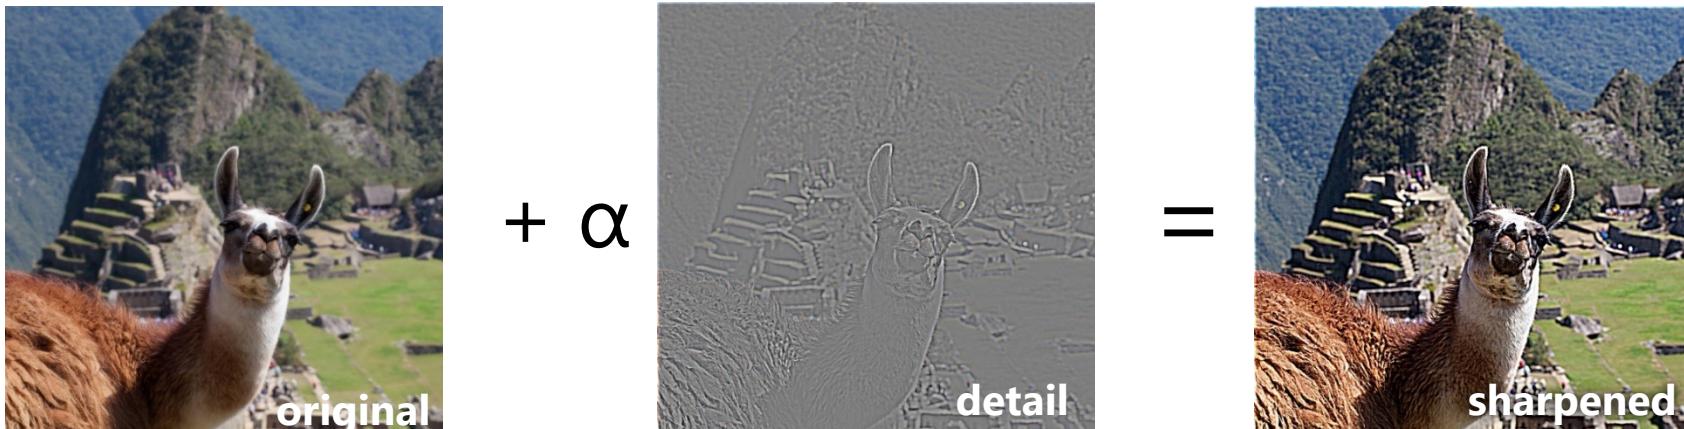
# Sharpening revisited

- What does blurring take away?



(This “detail extraction” operation is also called a **high-pass filter**)

Let's add it back:



# Sharpen filter



# Filters: Thresholding

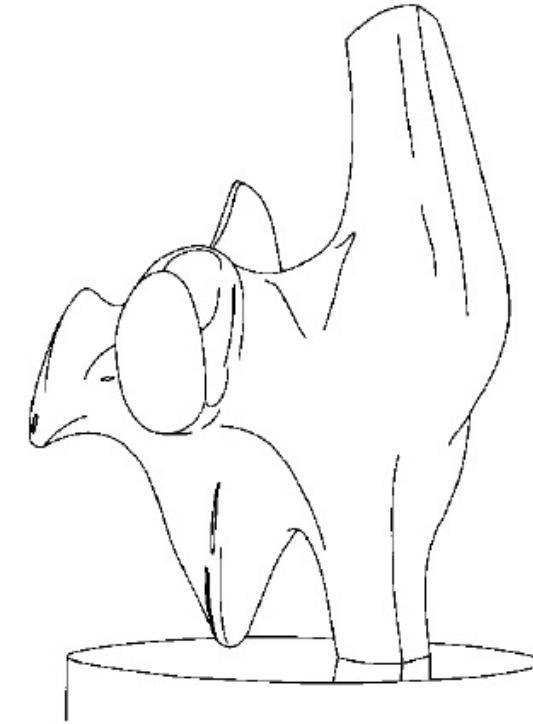
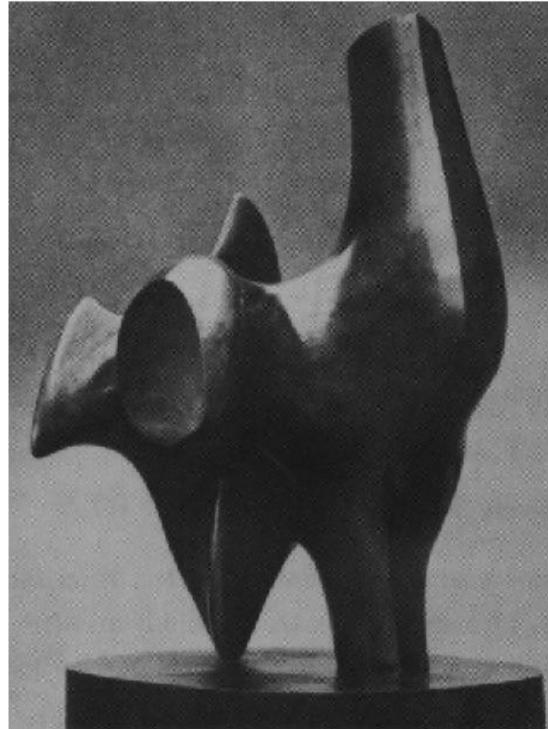


$$g(m, n) = \begin{cases} 255, & f(m, n) > A \\ 0 & otherwise \end{cases}$$

# Edge Detection

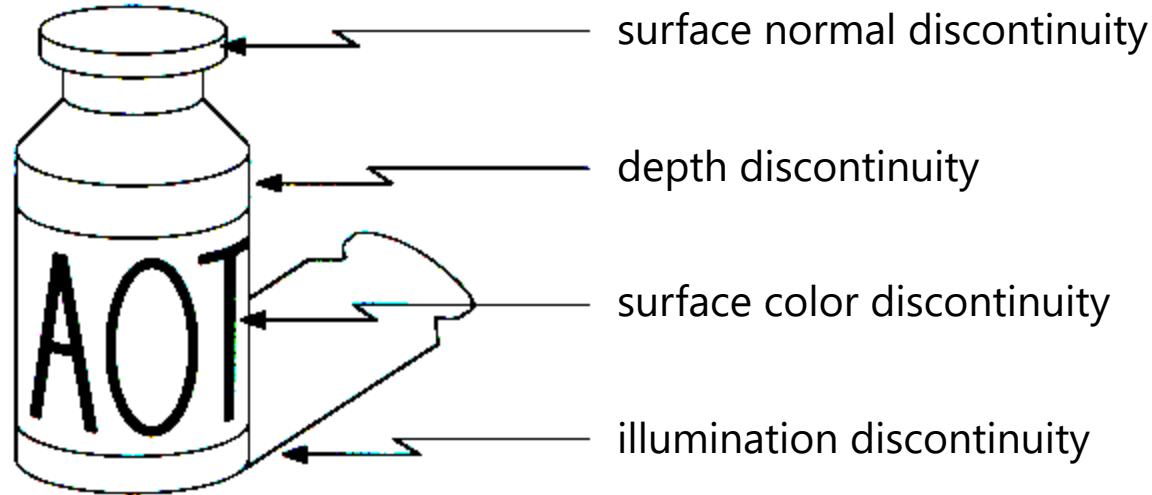
- Edge detection is an image processing technique for finding the boundaries of objects within images.
- It looks for discontinuities in brightness.

# Edge detection



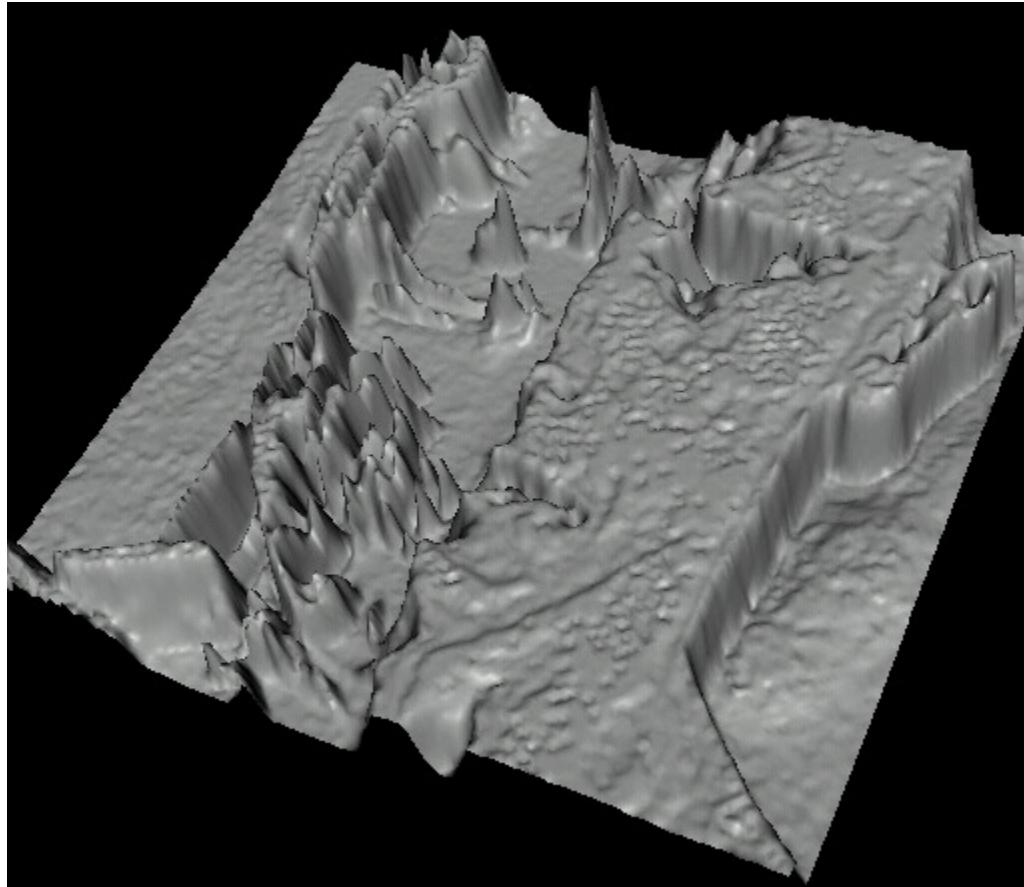
- Convert a 2D image into a set of curves
  - Extracts salient features of the scene
  - More compact than pixels

# Origin of edges



- Edges are caused by a variety of factors

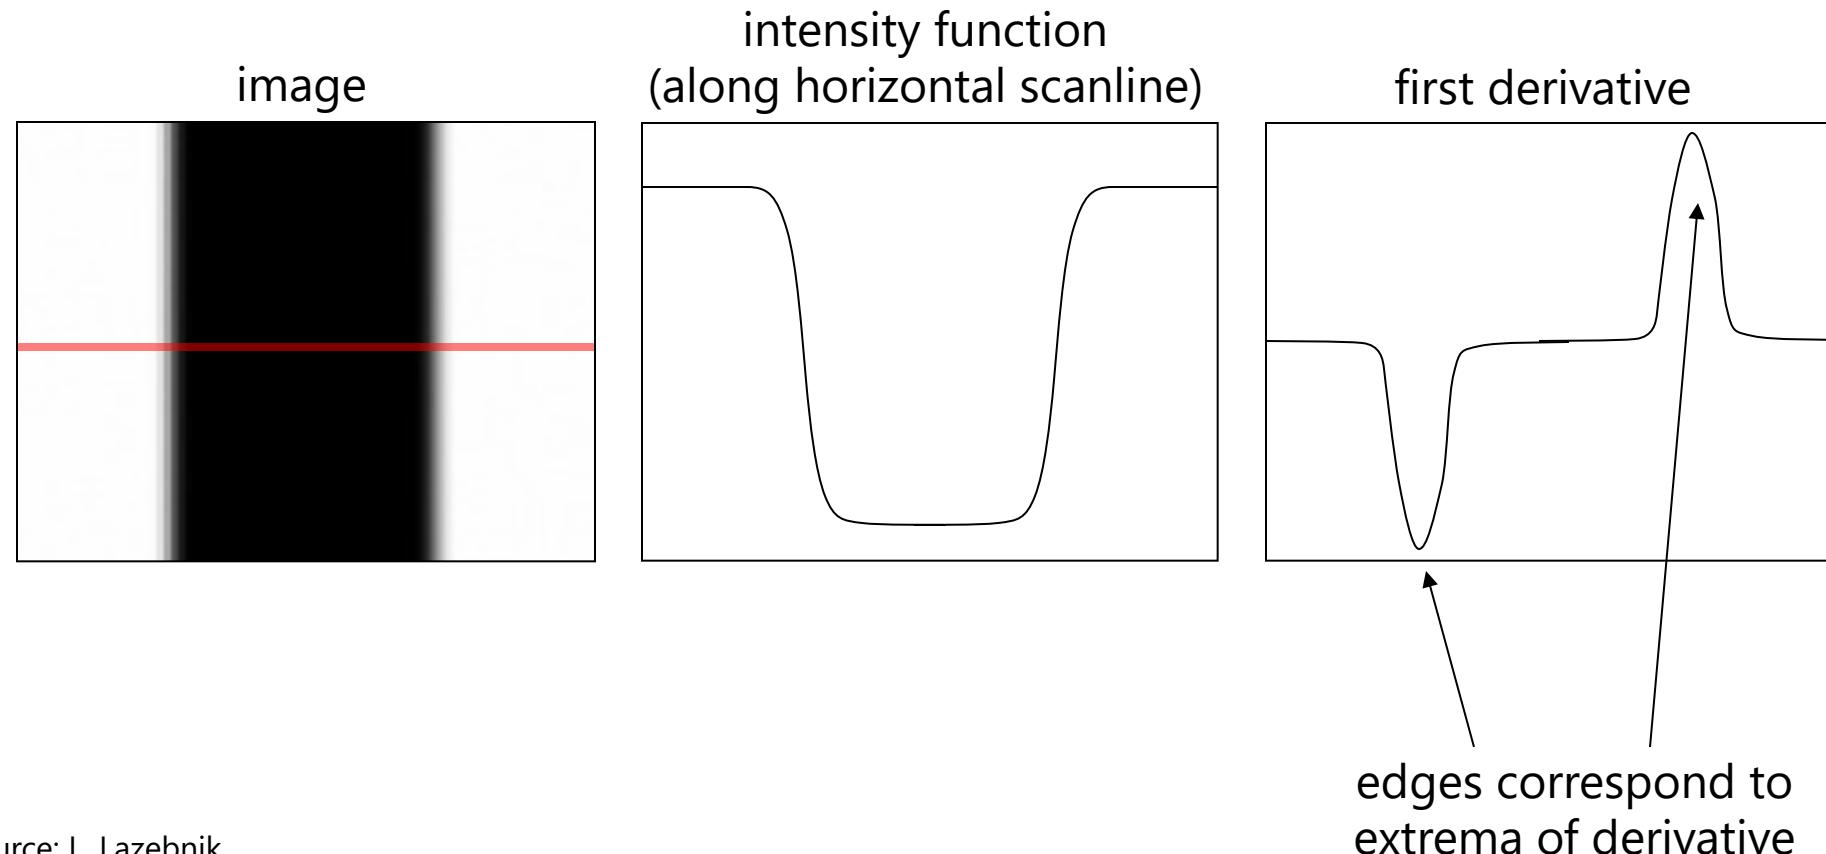
# Images as functions...



- Edges look like steep cliffs

# Characterizing edges

- An edge is a place of *rapid change* in the image intensity function

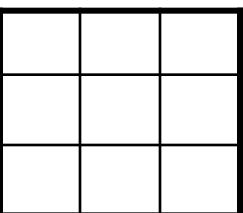


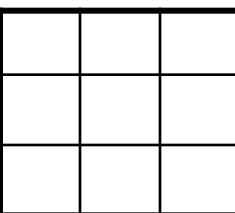
# Image derivatives

- How can we differentiate a *digital* image  $F[x,y]$ ?
  - Option 1: reconstruct a continuous image,  $f$ , then compute the derivative
  - Option 2: take discrete derivative (finite difference)

$$\frac{\partial f}{\partial x}[x, y] \approx F[x + 1, y] - F[x, y]$$

How would you implement this as a linear filter?

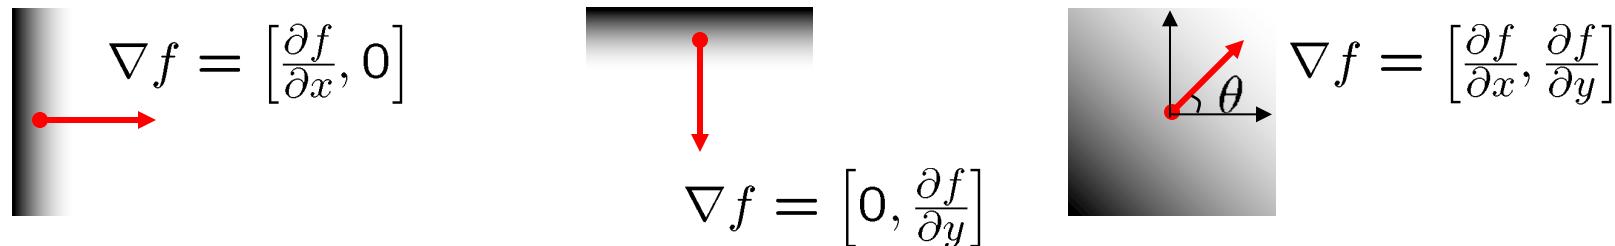
$$\frac{\partial f}{\partial x}:$$

$$H_x$$

$$\frac{\partial f}{\partial y}:$$

$$H_y$$

# Image gradient

- The *gradient* of an image:  $\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$

The gradient points in the direction of most rapid increase in intensity



The *edge strength* is given by the gradient magnitude:

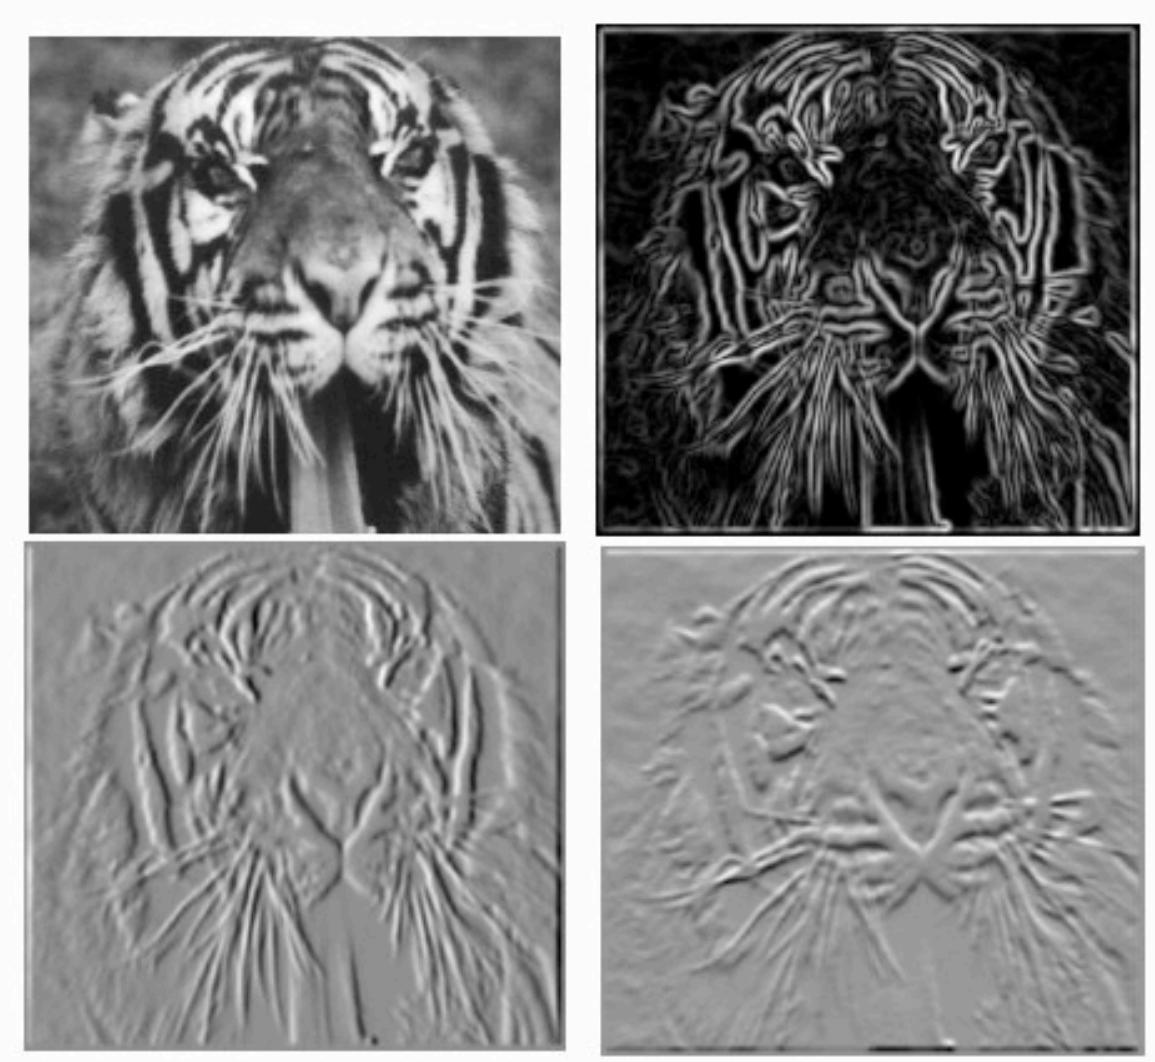
$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

The gradient direction is given by:

$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

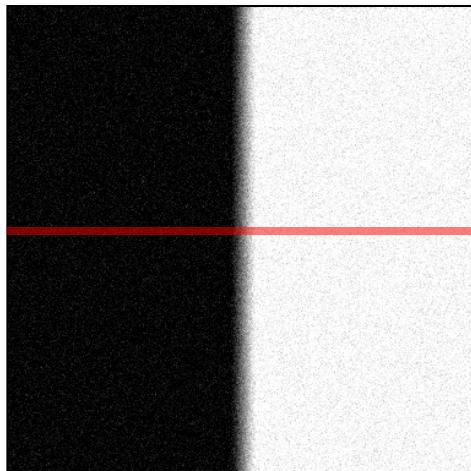
- how does this relate to the direction of the edge?

# Image gradient



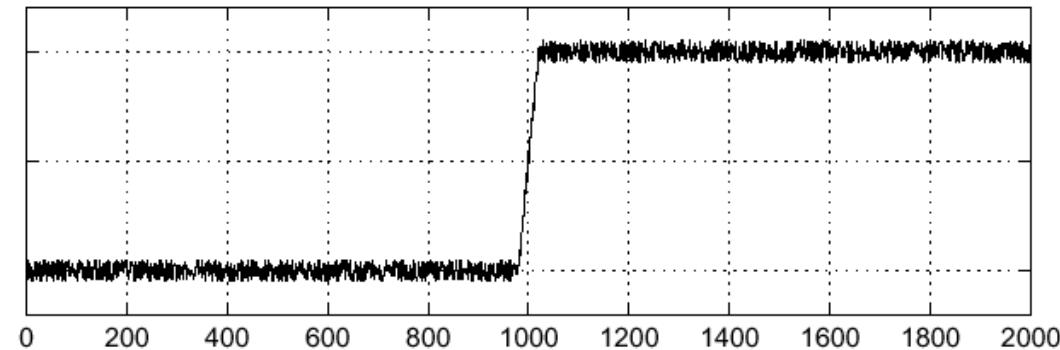
Source: L. Lazebnik

# Effects of noise

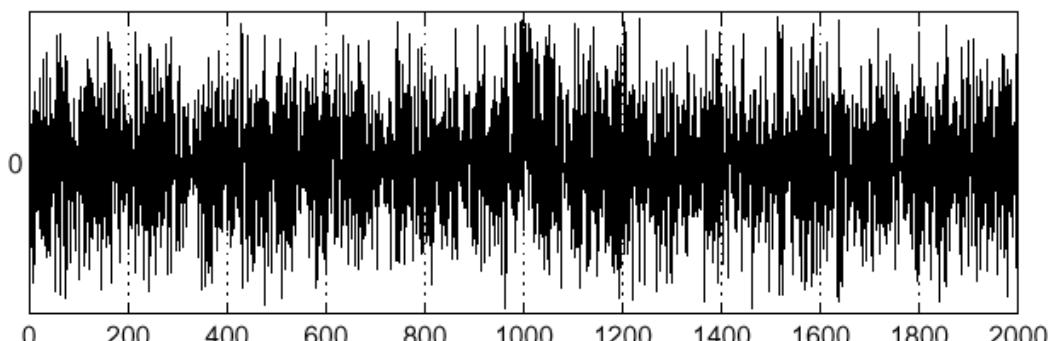


Noisy input image

$$f(x)$$



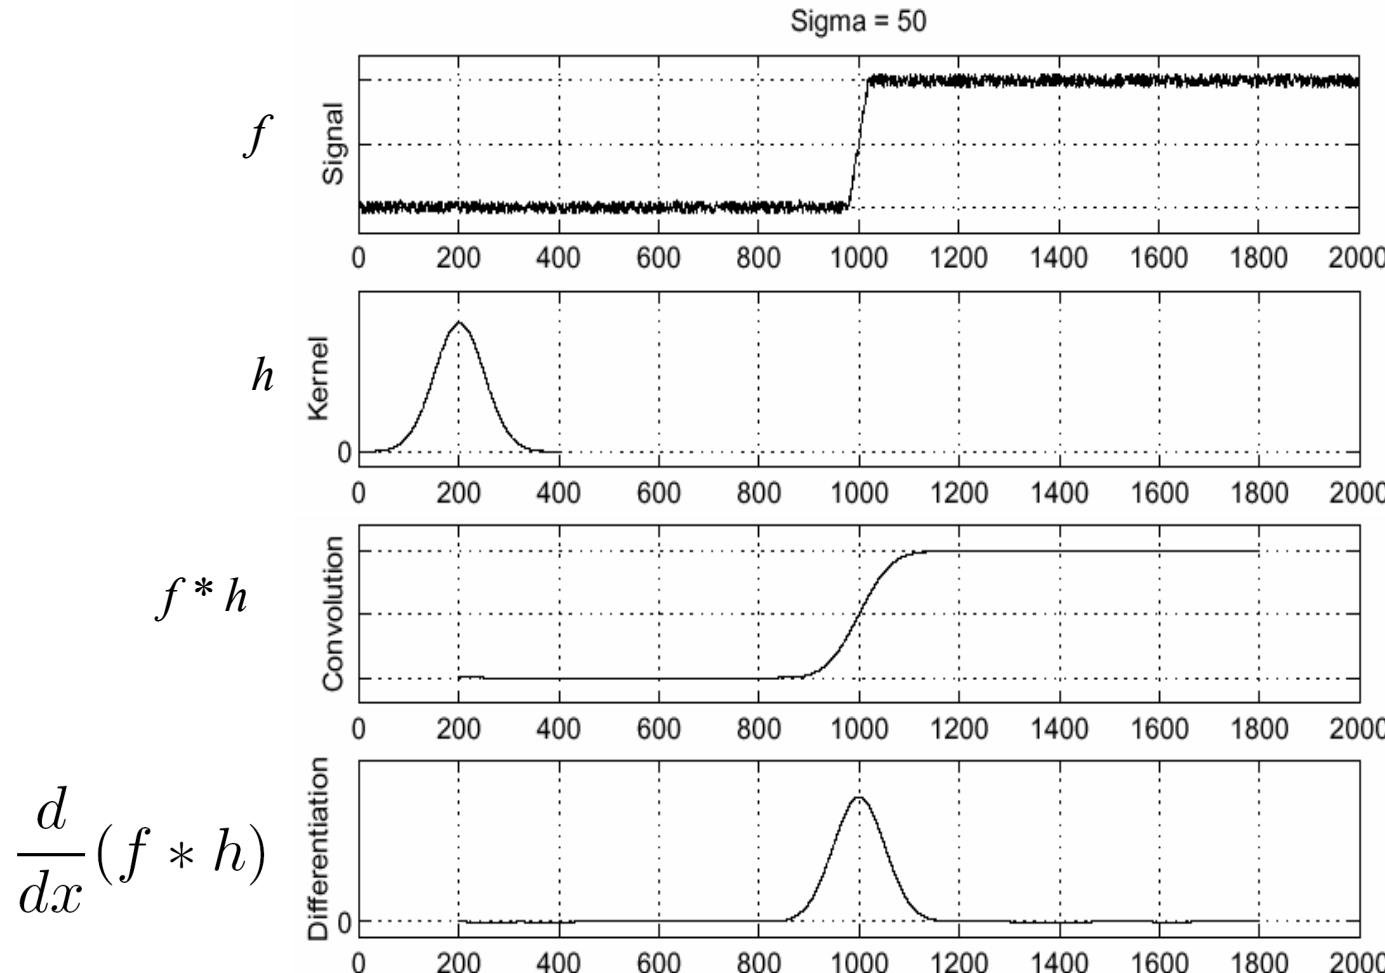
$$\frac{d}{dx}f(x)$$



Where is the edge?

Source: S. Seitz

# Solution: smooth first



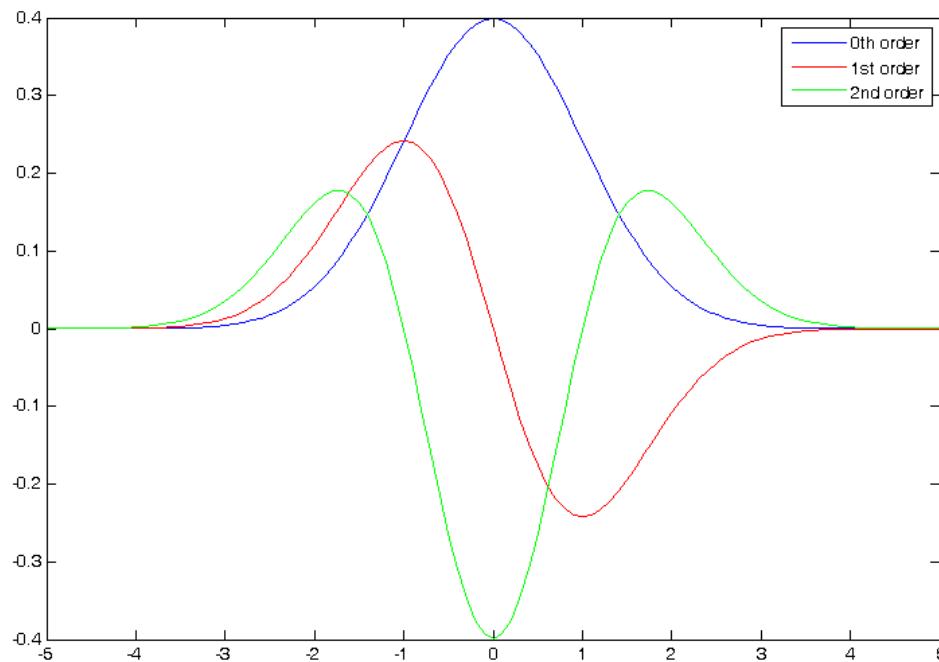
To find edges, look for peaks in  $\frac{d}{dx}(f * h)$

Source: S. Seitz

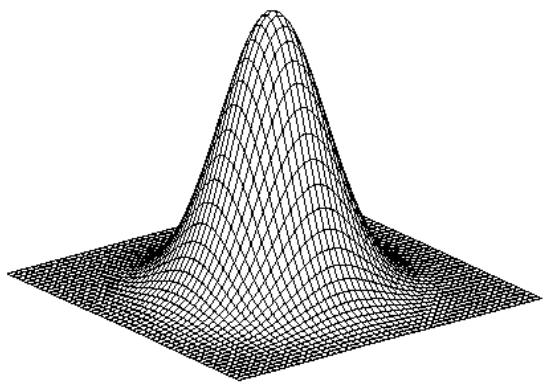
# The 1D Gaussian and its derivatives

$$G_\sigma(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}}$$

$$G'_\sigma(x) = \frac{d}{dx} G_\sigma(x) = -\frac{1}{\sigma} \left(\frac{x}{\sigma}\right) G_\sigma(x)$$

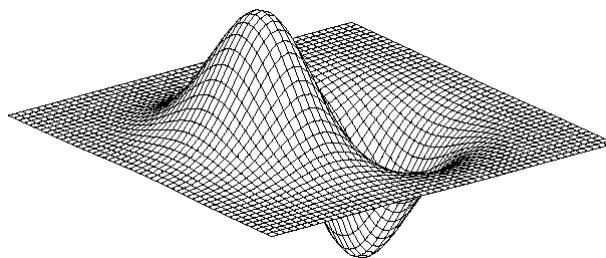


# 2D edge detection filters



Gaussian

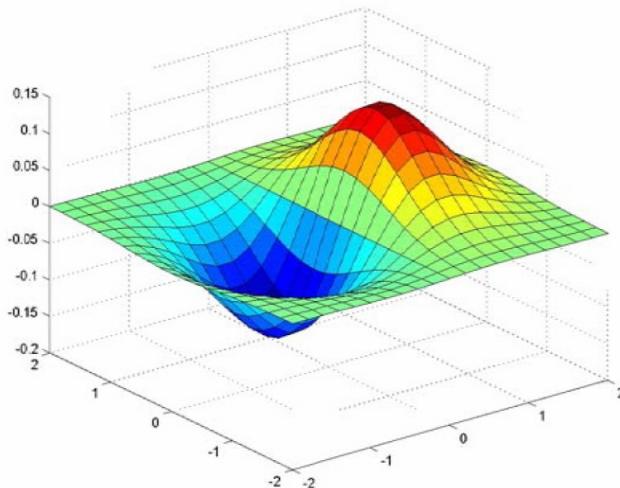
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



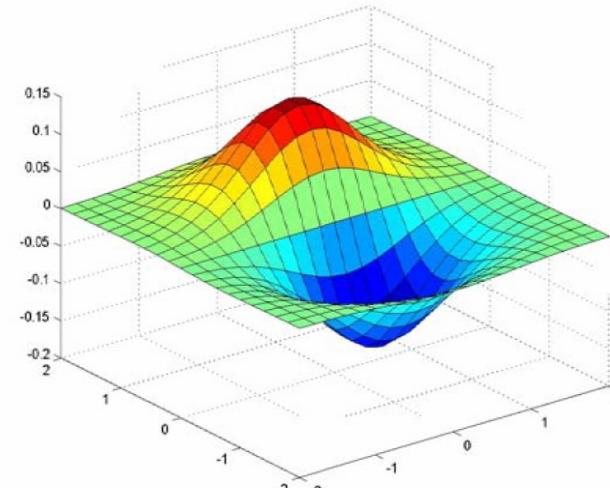
derivative of Gaussian ( $x$ )

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

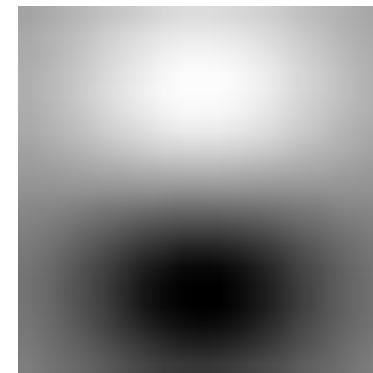
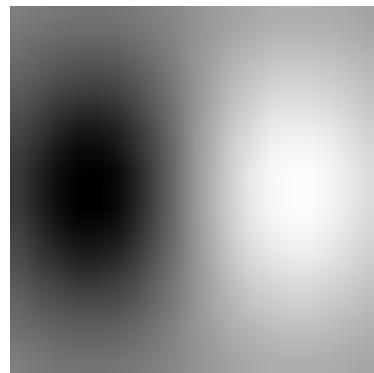
# Derivative of Gaussian filter



$x$ -direction



$y$ -direction



# The Sobel operator

- Common approximation of derivative of Gaussian

$$\frac{1}{8} \begin{matrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{matrix}$$

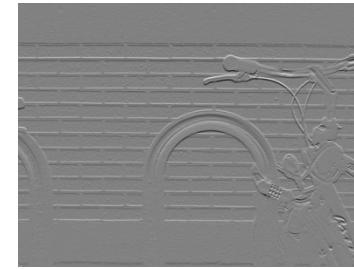
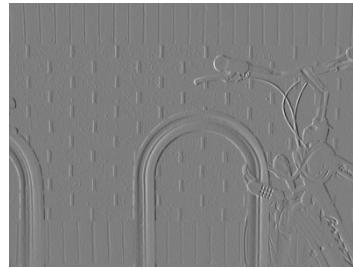
$s_x$

$$\frac{1}{8} \begin{matrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{matrix}$$

$s_y$

- The standard definition of the Sobel operator omits the 1/8 term
  - doesn't make a difference for edge detection
  - the 1/8 term **is** needed to get the right gradient magnitude

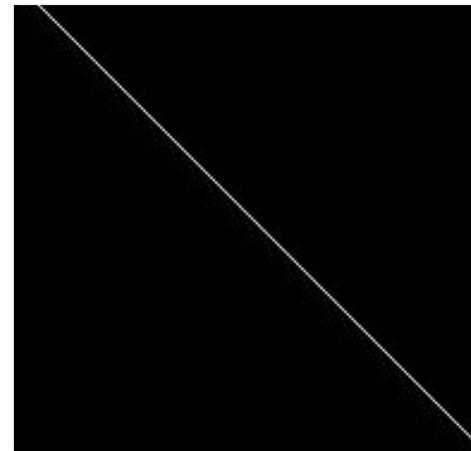
# Sobel operator: example



Source: Wikipedia



Image with Edge



Edge Location

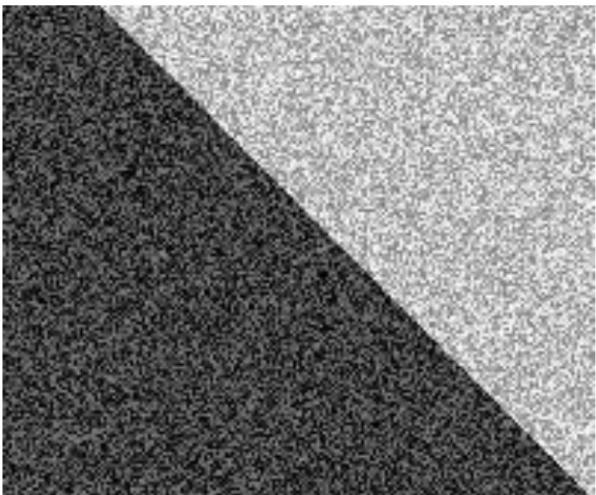
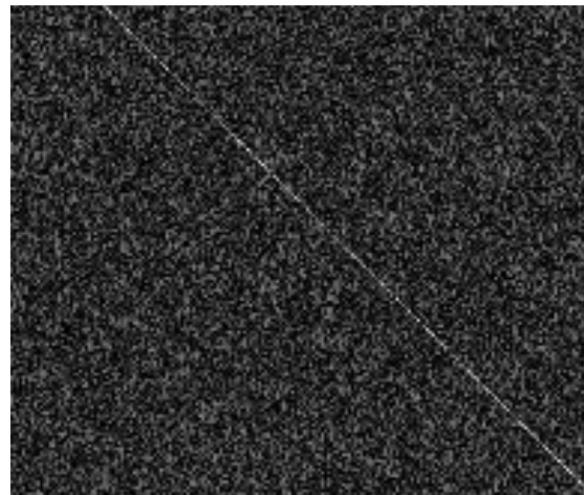
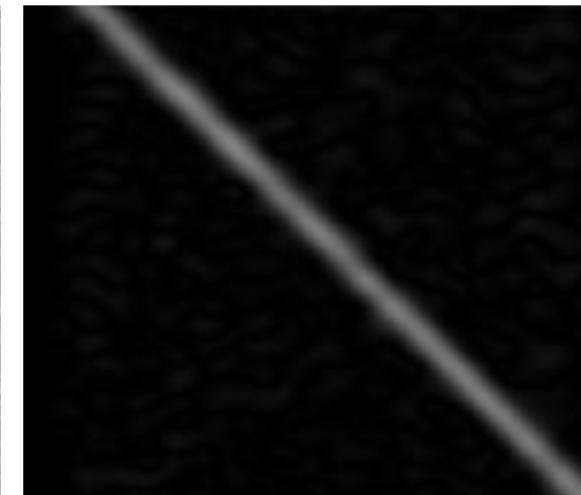


Image + Noise



Derivatives detect  
edge *and* noise



Smoothed derivative removes  
noise, but blurs edge

# Example



original image

Demo: <http://bigwww.epfl.ch/demo/ip/demos/edgeDetector/>

Image credit: Joseph Redmon

# Finding edges



smoothed gradient magnitude

# Finding edges

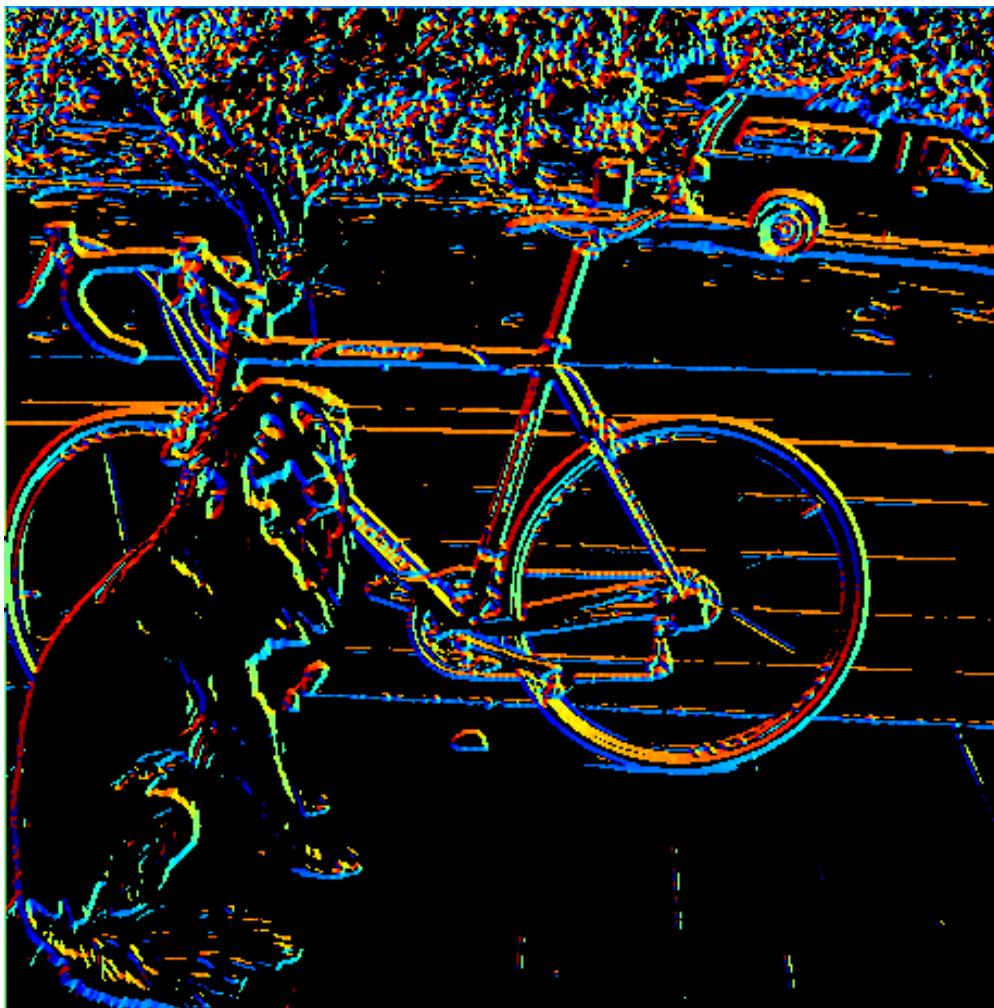


where is the edge?

thresholding

# Get Orientation at Each Pixel

- Get orientation (below, threshold at minimum gradient magnitude)



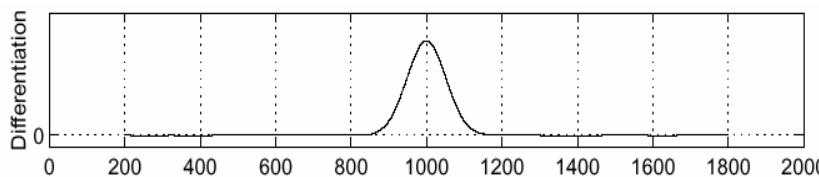
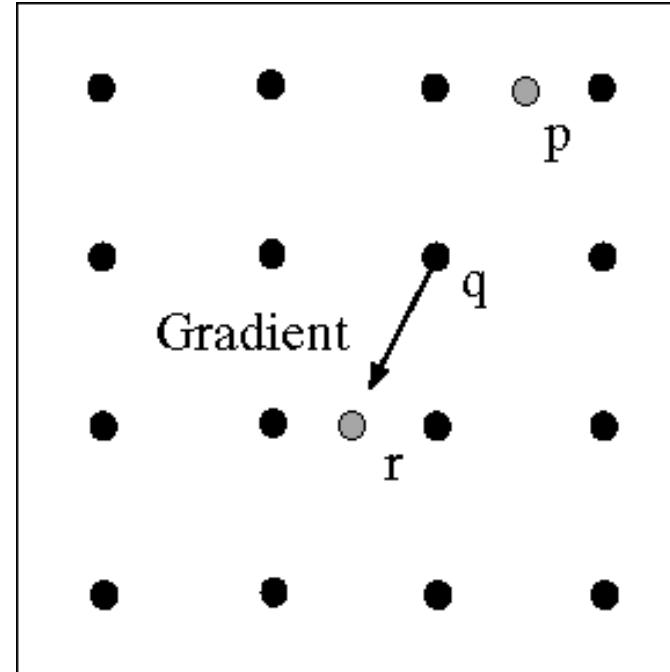
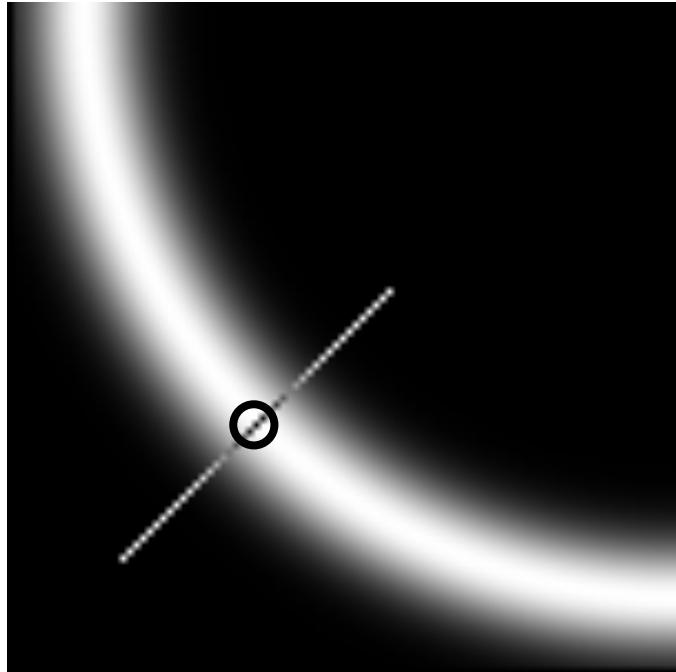
$\theta = \text{atan2}(gy, gx)$

360

**Gradient orientation angle**

0

# Non-maximum suppression



- Check if pixel is local maximum along gradient direction
  - requires *interpolating* pixels p and r

# Before Non-max Suppression



# After Non-max Suppression



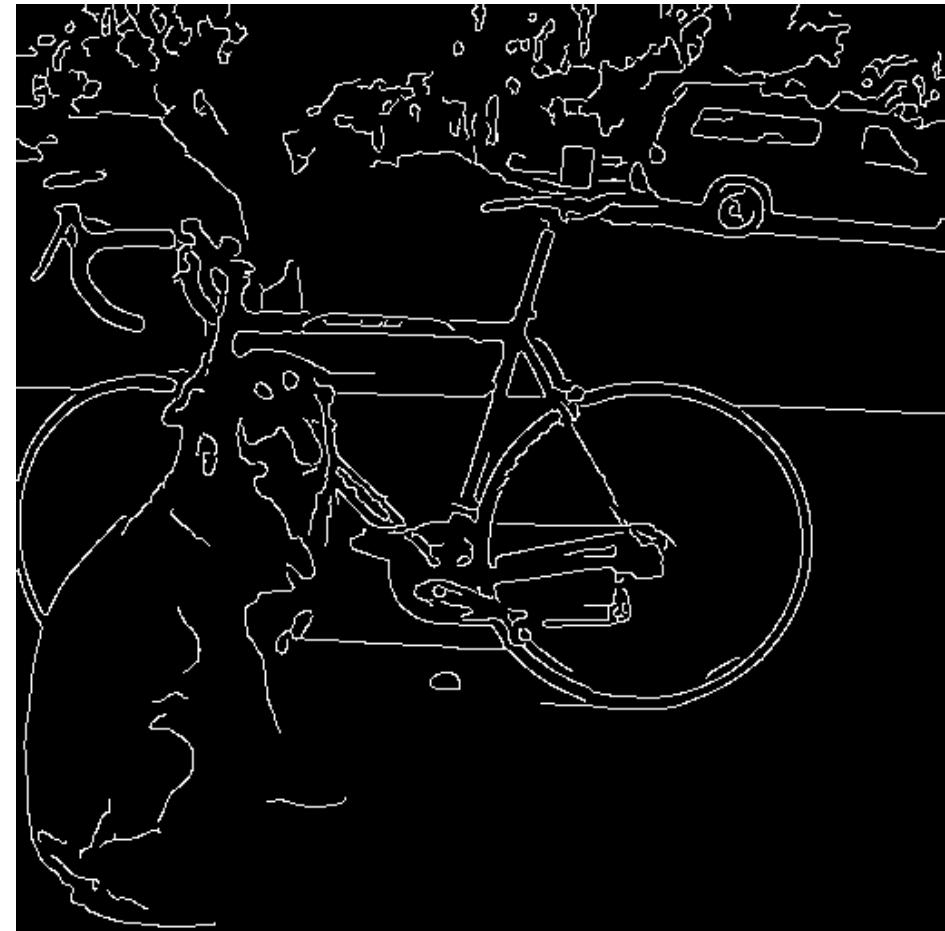
# Thresholding edges

- Still some noise
- Only want strong edges
- 2 thresholds, 3 cases
  - $R > T$ : strong edge
  - $R < T$  but  $R > t$ : weak edge
  - $R < t$ : no edge
- Why two thresholds?



# Connecting edges

- Strong edges are edges!
- Weak edges are edges if they connect to strong
- Look in some neighborhood (usually 8 closest)



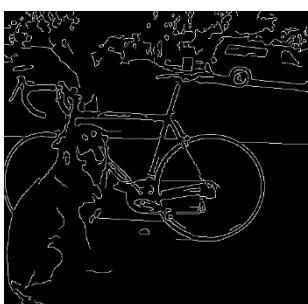
# Canny edge detector



MATLAB: `edge(image, 'canny')`



1. Filter image with derivative of Gaussian
2. Find magnitude and orientation of gradient
3. Non-maximum suppression
4. Linking and thresholding (hysteresis):
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them



# Canny edge detector

- Our first computer vision pipeline!
- Still a widely used edge detector in computer vision

J. Canny, [\*\*\*A Computational Approach To Edge Detection\*\*\*](#), IEEE Trans.  
Pattern Analysis and Machine Intelligence, 8:679-714, 1986.

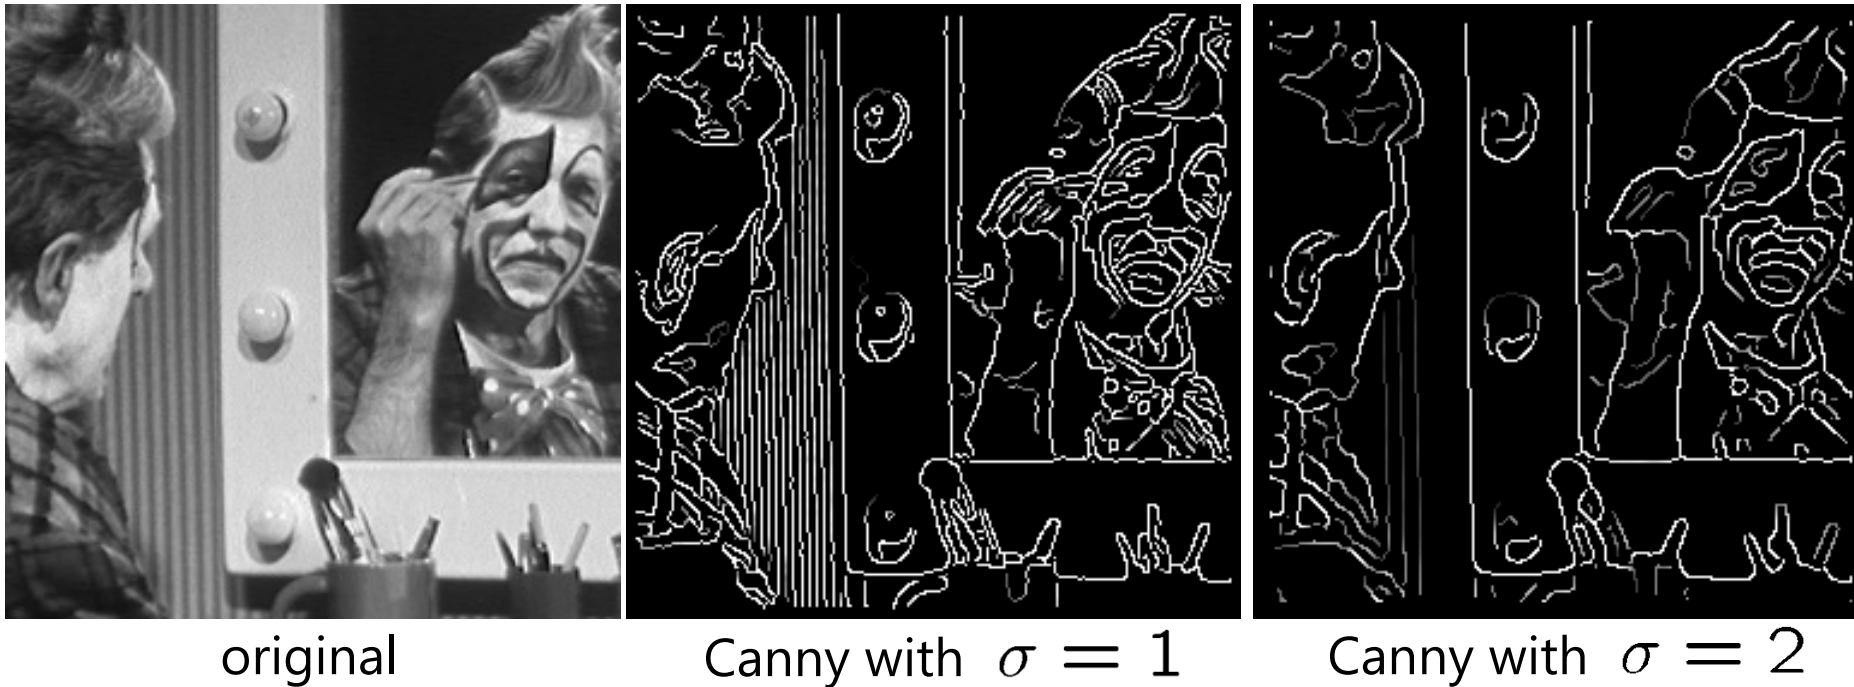
- Depends on several parameters:

high threshold

low threshold

$\sigma$  : width of the Gaussian blur

# Canny edge detector



- The choice of  $\sigma$  depends on desired behavior
  - large  $\sigma$  detects “large-scale” edges
  - small  $\sigma$  detects fine edges