


CSE 6367: Computer Vision

Principal Component Analysis

Slide Courtesy: Dr. Vassilis Athitsos,
University of Texas at Arlington

Vector Spaces

- For our purposes, a vector is a tuple of d numbers: $X = (x_1, x_2, \dots, x_d)$.
- An example vector space: the 2D plane.
 - Every point in the plot is a vector.
 - Specified by two numbers.



Images are Vectors

- An M -row \times N -column image is a vector of dimensionality ...

Images are Vectors

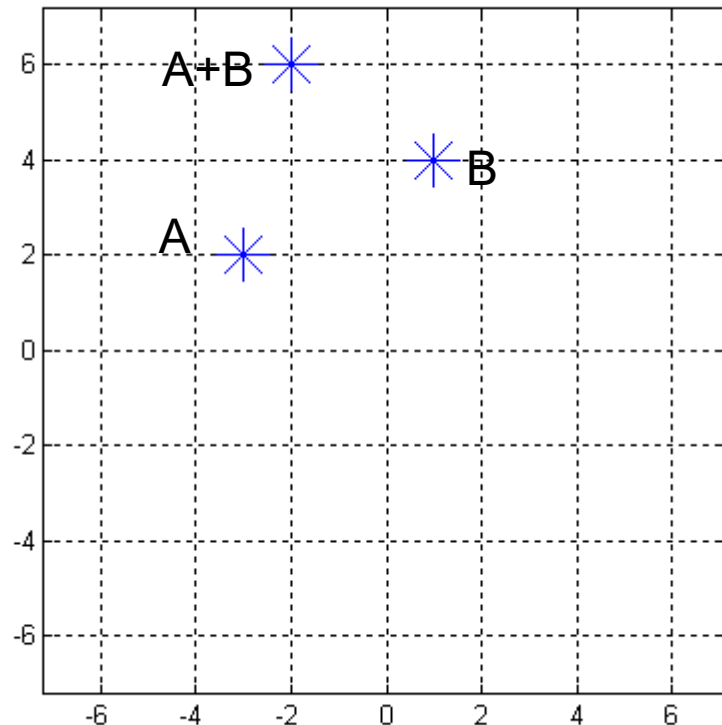
- An M -row x N -column image is a vector of dimensionality.
 - $M*N$ if the image is grayscale.
 - $M*N*3$ if the image is in color.

Images are Vectors

- Consider a 4x3 grayscale image:
 - $A = \begin{bmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \\ A_{41} & A_{42} & A_{43} \end{bmatrix};$
- The (Matlab) vector representation A_v of A is:
 $A_v = [A_{11} \ A_{21} \ A_{31} \ A_{41} \ A_{12} \ A_{22} \ A_{32} \ A_{42} \ A_{13} \ A_{23} \ A_{33} \ A_{43}];$
- Mathematically, order does not matter, IF IT IS THE SAME FOR ALL IMAGES.
- In Matlab, to vectorize an image:
 - $A_v = A(:);$
 - NOTE: The above returns a COLUMN vector.


Vector Operations: Addition

- $(x_1, \dots, x_d) + (y_1, \dots, y_d) = (x_1+y_1, \dots, x_d+y_d)$
- Example: in 2D:
 - $A = (-3, 2)$
 - $B = (1, 4)$
 - $A+B = (-2, 6)$.




Vector Operations: Addition


- $(x_1, \dots, x_d) + (y_1, \dots, y_d) = (x_1+y_1, \dots, x_d+y_d)$
- Example: in 2D:
 - $A = (-3, 2)$
 - $B = (1, 4)$
 - $A+B = (-2, 6)$.



Addition in Image Space



+



=

MxN image
all pixels (255, 0, 0)

MxN image
all pixels (0, 255, 0)




+




=


Addition in Image Space



+



=



MxN image
all pixels (255, 0, 0)

MxN image
all pixels (0, 255, 0)

MxN image
all pixels (255, 255, 0)



+




=




What is the range
of pixel values here?


Addition in Image Space



+



=



MxN image
all pixels (255, 0, 0)

MxN image
all pixels (0, 255, 0)

MxN image
all pixels (255, 255, 0)



+



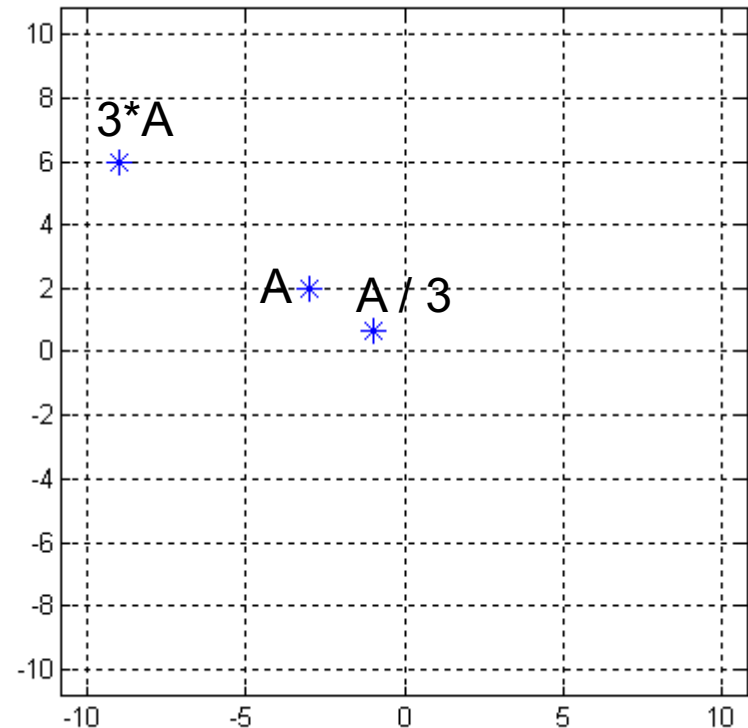
=



Range of pixel values:
from 0 to 2*255

Vector Operations: Scalar Multiplication

- $c * (x_1, \dots, x_d) = (c * x_1, \dots, c * x_d)$
- $(x_1, \dots, x_d) / c = 1/c * (x_1, \dots, x_d)$
 $= (x_1/c, \dots, x_d/c)$
- Example: in 2D:
 - $A = (-3, 2)$
 - $3 * A = (-9, 6)$
 - $A / 3 = (-1, 0.67)$.



Multiplication in Image Space

image



image * 0.7



image * 0.5

Operations in Image Space

- Note: with addition and multiplication we can easily generate images with values outside the $[0, 255]$ range.
 - These operations are perfectly legal.
 - However, we cannot visualize the results directly.
 - Must convert back to $[0, 255]$ range to visualize.

Linear Combinations

- Example: $c_1 * v_1 + c_2 * v_2 + c_3 * v_3$
 - c_1, c_2, c_3 : real numbers.
 - v_1, v_2, v_3 : vectors
 - result:

Linear Combinations


- Example: $c_1 * v_1 + c_2 * v_2 + c_3 * v_3$
 - c_1, c_2, c_3 : real numbers.
 - v_1, v_2, v_3 : vectors
 - result: vector.

Vectors Must Be of Same Size

- We cannot add vectors that are not of the same size.
 - $(1,2) + (0,1,0)$ is NOT DEFINED.
- To add images A and B:
 - IMPORTANT: Most of the time, it only makes sense to add A and B ONLY if they have the same number of rows and the same number of columns.
 - WARNING: Matlab will happily do the following:

```
a = rand(4,3);  
b = rand(6,2);  
c = a(:) + b(:);
```


Example Linear Combination



a

b

c

d

e



avg


```
a = read_gray('4309d111.bmp');  
b = read_gray('4286d201.bmp');  
c = read_gray('4846d101.bmp');  
d = read_gray('4848d101.bmp');  
e = read_gray('4853d101.bmp');  
avg = 0.2*a + 0.2*b + 0.2*c + 0.2*d + 0.2*e;
```

```
% or, equivalently:
```

```
avg = (a+b+c+d+e) / 5;
```


Dimensionality Reduction

- Consider a set of vectors in a d -dimensional space.
- How many numbers do we need to represent each vector?




Dimensionality Reduction

- Consider a set of vectors in a d -dimensional space.
- How many numbers do we need to represent each vector?
 - At most d : the same as the number of dimensions.
- Can we use fewer?




Dimensionality Reduction

- Consider a set of vectors in a d -dimensional space.
- How many numbers do we need to represent each vector?
 - At most d : the same as the number of dimensions.
- Can we use fewer?




Dimensionality Reduction

- In this example, every point (x, y) is on a line
– $y = ax + b$;
- If we have 100 points on this plot, how many numbers do we need to specify them?




Dimensionality Reduction

- In this example, every point (x, y) is on a line
 - $y = ax + b$;
- If we have 100 points on this plot, how many numbers do we need to specify them?
 - 102: a , b , and the x coordinate of each point.
 - Asymptotically: one number per point.




Lossy Dimensionality Reduction

- Suppose we want to project all points to a single line.
- This will be *lossy*.
- What would be the best line?




Lossy Dimensionality Reduction

- Suppose we want to project all points to a single line.
- This will be *lossy*.
- What would be the best line?
- Optimization problem.
 - Infinite answers.
 - We must define how to evaluate each answer.




Optimization Criterion

- We want to measure how good a projection P is, GIVEN A SET OF POINTS.
 - If we don't have a specific set of points in mind, what would be the best projection?




Optimization Criterion

- We want to measure how good a projection P is, GIVEN A SET OF POINTS.
 - If we don't have a specific set of points in mind, what would be the best projection?
 - NONE: all are equally good/bad.




Optimization Criterion

- Consider a pair of points: X_1, X_2 .
- $D1$ = squared distance from X_1 to X_2 .
 - $\text{sum}((X_1 - X_2) .* (X_1 - X_2))$
- $D2$ = squared distance from $P(X_1)$ to $P(X_2)$.
- $\text{Error}(X_1, X_2) = D1 - D2$.
 - Will it ever be negative?




Optimization Criterion

- Consider a pair of points: X_1, X_2 .
- D_1 = squared distance from X_1 to X_2 .
 - $\text{sum}((X_1 - X_2) .* (X_1 - X_2))$
- D_2 = squared distance from $P(X_1)$ to $P(X_2)$.
- $\text{Error}(X_1, X_2) = D_1 - D_2$.
 - Will it ever be negative?
 - NO: $D_1 \geq D_2$ always.



Optimization Criterion


- Now, consider the entire set of points:
 - X_1, X_2, \dots, X_n .
- $\text{Error}(P) = \sum (\text{Error}(X_i, X_j) \mid i, j = 1, \dots, n, i \neq j)$.



- Interpretation:
 - We measure how well P preserves distances.
 - If P preserves distances, $\text{Error}(P) = 0$.

Example: Perfect Projection Exists

- In this case, projecting to a line oriented at 30 degrees would give zero error.




Finding the Best Projection: PCA

- First step: center the data.


```
number = size(points, 2);  
% note that we are transposing twice  
average = [mean(points')]';  
centered_points = zeros(size(points));  
  
for index = 1:number  
    centered_points(:, index) = points(:, index) - average;  
end  
  
plot_points(centered_points, 2);
```

Finding the Best Projection: PCA

- First step: center the data.



points



centered_points

Finding the Best Projection: PCA

- Second step: compute the covariance matrix.

```
covariance_matrix = centered_points * centered_points';
```

- In the above line we assume that each column is a vector.

Finding the Best Projection: PCA

- Second step: compute the covariance matrix.

```
covariance_matrix = centered_points * centered_points';
```

- In the above line we assume that each column is a vector.
- Third step: compute the eigenvectors and eigenvalues of the covariance matrix.

```
[eigenvectors eigenvalues] = eig(covariance_matrix);
```

Eigenvectors and Eigenvalues

```
eigenvectors =  
    0.4837    -0.8753  
   -0.8753    -0.4837
```


```
eigenvalues =  
    2.0217         0  
    0    77.2183
```

- Each eigenvector v is a column, that specifies a line going through the origin.
- The importance of the i -th eigenvector is reflected by the i -th eigenvalue.
 - second eigenvalue = 77, first eigenvalue = 2, => second eigenvector is far more important.

Visualizing the Eigenvectors

black: v_1 (eigenvalue = 2.02)

red: v_2 (eigenvalue = 77.2)



```
plot_points(points, 1);  
p1 = eigenvectors(:, 1);  
p2 = eigenvectors(:, 2);  
plot([0, p1(1)], [0, p1(2)], 'k-', 'linewidth', 3);  
hold on;  
plot([0, p2(1)], [0, p2(2)], 'r-', 'linewidth', 3);
```

PCA Code

```
function [average, eigenvectors, eigenvalues] = ...
    compute_pca(vectors)

number = size(vectors, 2);
% note that we are transposing twice
average = [mean(vectors')]';
centered_vectors = zeros(size(vectors));

for index = 1:number
    centered_vectors(:, index) = vectors(:, index) - average;
end

covariance_matrix = centered_vectors * centered_vectors';
[eigenvectors eigenvalues] = eig(covariance_matrix);

% eigenvalues is a matrix, but only the diagonal
% matters, so we throw away the rest
eigenvalues = diag(eigenvalues);
[eigenvalues, indices] = sort(eigenvalues, 'descend');
eigenvectors = eigenvectors(:, indices);
```

Reducing Dimensions From 2 to 1


- Precompute:
 - the eigenvector P_1 with the largest eigenvalue.
 - the mean **avg** of all the data.

Reducing Dimensions From 2 to 1

- Two key operations:
 - projection to lower-dimensional space.
 - backprojection to the original space.
- Projection:
 - $P(V) = \langle V - \text{avg}, P1 \rangle = P1' * (V - \text{avg})$
 - Dot product between $(V - \text{avg})$ and $P1$.
 - NOTE: The eigenvectors that Matlab returns have unit norm.
 - Backprojection:
 - $B(P(V)) = P1 * P(V) + \text{avg}$.


Example: From 2 Dimensions to 1

- A set of points.




Example: From 2 Dimensions to 1

- Do PCA, identify p_1 = the first eigenvector.
- We plot the line of direction of p_1 .




Example: From 2 Dimensions to 1

- Choose a point $v4 = [-1.556, 0.576]'$.
 - Shown in red.




Example: From 2 Dimensions to 1

- $\text{centered_v4} = \text{v4} - \text{average}$.
 - Shown in cyan.




Example: From 2 Dimensions to 1

- $\text{projection} = p1' * (\text{centered_v4})$;
- result: projection = 1.43 (shown in magenta)




Example: From 2 Dimensions to 1

- Note: projection is a single number. To plot it, we actually treat that number as the x value, and use a y value of 0.




Example: From 1 Dimension to 2

- $b_1 = p_1 * \text{projection}$;
 - shown in red, on top of black line.



Example: From 1 Dimension to 2

- reconstructed $v_4 = b_1 + \text{average}$;
– shown in green.



PCA on Faces

- Motivation: If a face is a 31×25 window, we need 775 numbers to describe the face.
- With PCA, we can approximate face images with much fewer numbers.
- Two benefits:
 - Faster computations.
 - The amount of loss can be used for face detection.

PCA vs Template Matching

- If we use template matching to detect faces, what is the perfect face (easiest to be detected)?
- How about PCA?

PCA vs Template Matching

- Template matching:
 - The average face is the only perfect face (after we normalize for brightness and contrast).
 - As a model, it is not rich enough.
- PCA:
 - There is a space of perfect faces.

Preparing a Dataset

- Get a set of face images.
 - They must be of same size, and aligned, so that eye locations match each other, nose locations match each other, and so on.
 - Make the mean and std of all faces equal.
 - Discards variations in intensity and contrast.

Code for Preprocessing Faces


```
clear;
load faces1032;

number = size(faces, 2);
dimensions = size(faces, 1);

% set mean of all faces to 0, and std to 1.
for index = 1: number
    face = faces(:, index);
    face = (face - mean(face)) / std(face);
    faces(:, index) = face;
end

% do pca
[mean_face, eigenvectors, eigenvalues] = compute_pca(faces);
```

Visualizing Eigenfaces




Approximating a Face With 0 Numbers

- What is the best approximation we can get for a face image, if we know nothing about the face image (except that it is a face)?



Equivalent Question in 2D


- What is our best guess for a 2D point from this point cloud, if we know nothing about that 2D point (except that it belongs to the cloud)?



Equivalent Question in 2D

- What is our best guess for a 2D point from this point cloud, if we know nothing about that 2D point (except that it belongs to the cloud)?

Answer: the average of all points in the cloud.



Approximating a Face With 0 Numbers

- What is the best approximation we can get for a face image, if we know nothing about the face image (except that it is a face)?




Approximating a Face With 0 Numbers

- What is the best approximation we can get for a face image, if we know nothing about the face image (except that it is a face)?
 - The average face.




Guessing a 2D Point Given 1 Number

- What is our best guess for a 2D point from this point cloud, if we know nothing about that 2D point, except a single number?
 - What should that number be?



Guessing a 2D Point Given 1 Number

- What is our best guess for a 2D point from this point cloud, if we know nothing about that 2D point, except a single number?
 - What should that number be?
 - Answer: the projection on the first eigenvector.



Approximating a Face With 1 Numbers

- What is the best approximation we can get for a face image, if we can represent the face with a single number?



Approximating a Face With 1 Numbers

- With 0 numbers, we get the average face.
- 1 number: PCA projection to 1D.
- Reconstruction:
 - `average face + <face, eigenvector1> * eigenvector1`



Approximating a Face With 1 Number

```
v1 = faces(:, 13);  
f1 = reshape(v1, size(mean_face));  
k = 0;  
%%  
x1 = v1' * eigenvectors(:, 1);  
term1 = x1 * eigenvectors(:, 1);  
term1 = reshape(term1, size(mean_face));  
  
reconstructed1 = mean_face + term1;
```



mean_face

Approximating a Face With 1 Number

```
v1 = faces(:, 13);  
f1 = reshape(v1, size(mean_face));  
k = 0;  
%%  
x1 = v1' * eigenvectors(:, 1);  
term1 = x1 * eigenvectors(:, 1);  
term1 = reshape(term1, size(mean_face));  
  
reconstructed1 = mean_face + term1;
```



reconstructed1

Approximating a Face With 2 Numbers

```
x2 = v1' * eigenvectors(:, 2);  
term2 = x2 * eigenvectors(:, 2);  
term2 = reshape(term2, size(mean_face));  
  
reconstructed2 = reconstructed1 + term2;
```



reconstructed2

Approximating a Face With 3 Numbers

```
x = v1' * eigenvectors(:, 3);  
term = x * eigenvectors(:, 3);  
term = reshape(term1, size(mean_face));  
  
reconstructed3 = reconstructed2 + term;
```



reconstructed3

Approximating a Face With 4 Numbers



reconstructed4

Approximating a Face With 5 Numbers



reconstructed5

Approximating a Face With 6 Numbers



reconstructed6

Approximating a Face With 7 Numbers



reconstructed7

Approximating a Face With 10 Numbers

```
f1 = faces(:, 13);  
f1 = reshape(f1, size(mean_face));  
  
ev = eigenvectors(:, 1:10);  
p = pca_projection(f1, mean_face, ev);  
b = pca_backprojection(p, mean_face, ev);  
reconstructed10 = reshape(b, size(mean_face));
```



reconstructed10

PCA Projection Code

```
function result = normalize_face(image_window, mean_face)

% function result = normalize_face(vector)
%
% normalizes the vector so that the size matches that of \
% mean face, the mean is 0 and the std is 1.

result = imresize(image_window, size(mean_face), 'bilinear');
result = result(:);
result = result - mean(result(:));
result = result / std(result(:));
result(isnan(result)) = 0;
```


PCA Projection Code

```
function result = eigenface_projection(image_window, ...  
                                       mean_face, eigenvectors)  
  
% function result = eigenface_projection(image_window, ...  
%                                       average, eigenvectors)  
%  
% the vector is converted to a column vector. Each column in  
% eigenvectors should be an eigenvector. The mean is converted  
% to a column vector  
  
normalized = normalize_face(image_window, mean_face);  
% subtract mean from vector  
centered = normalized(:) - mean_face(:);  
  
% convert vector to a column vector.  
result = eigenvectors' * centered;
```

PCA Backprojection Code

```
function result = pca_backprojection(projection, ...  
                                     average, eigenvectors)  
  
projection = projection(:);  
centered_result = eigenvectors * projection;  
result = centered_result + average(:);
```

- The PCA projection gives us a few numbers to represent a face.
- The backprojection uses those few numbers to generate a face image.

Projection/Backprojection Results



original



10 eigenfaces



40 eigenfaces



100 eigenfaces

Projection/Backprojection Results



original



10 eigenfaces



40 eigenfaces



100 eigenfaces

Projection/Backprojection Results



original



10 eigenfaces



40 eigenfaces



100 eigenfaces

Note: teeth not visible using 10 eigenfaces

Projection/Backprojection Results



original



10 eigenfaces



40 eigenfaces



100 eigenfaces

Note: using 10 eigenfaces, gaze direction is towards camera

Projection/Backprojection Results



original



10 eigenfaces



40 eigenfaces



100 eigenfaces

Note: using 10 eigenfaces, glasses are removed

Projecting a Non-Face



original



10 eigenfaces



40 eigenfaces



100 eigenfaces

How Much Can 10 Numbers Tell?



original



eigenfaces

Using Eigenfaces for Detection

- Intuition: Why are eigenfaces good for representing faces?
 - Because projection OF FACES to the eigenspace loses little information.
 - However, projection of NON-FACES to the eigenspace of faces loses more information.
- Criterion for face detection: how much is the reconstruction error?
- Quantitative:
 - Define f = face, b = $\text{backprojection}(\text{projection}(f))$.
 - Error: norm of $(f - b)$.

Evaluating Error for a Window

```
function result = pca_score(window, mean_face, ...  
                             eigenvectors, eigenface_number)  
  
% normalize the window  
window = imresize(window, size(mean_face), 'bilinear');  
window = window(:);  
window = window - mean(window);  
window = window / std(window);  
window(isnan(window)) = 0;  
  
top_eigenvectors = eigenvectors(:, 1:eigenface_number);  
projection = pca_projection(window, mean_face, ..  
                             top_eigenvectors);  
reconstructed = pca_backprojection(projection, mean_face, ...  
                                   top_eigenvectors);  
  
diff = reconstructed - window(:);  
result = sum(diff .* diff);
```

Review: Correlation-based Detection

```
function [max_responses, max_scales, max_rotations] = ...  
    template_search(image, template, scales, rotations, result_number)  
  
% function [result, max_scales, max_rotations] = ...  
%     template_search(image, template, scales, rotations, result_number)  
%  
% for each pixel, search over the specified scales and rotations,  
% and record:  
% - in result, the max normalized correlation score for that pixel  
%   over all scales  
% - in max_scales, the scale that gave the max score  
% - in max_rotations, the rotation that gave the max score  
%  
% clockwise rotations are positive, counterclockwise rotations are  
% negative.  
% rotations are specified in degrees
```

Review: Correlation-based Detection

```
function [max_responses, max_scales, max_rotations] = ...  
    template_search(image, template, scales, rotations, result_number)
```

```
max_responses = ones(size(image)) * -10;  
max_scales = zeros(size(image));  
max_rotations = zeros(size(image));
```

```
for rotation = rotations  
    rotated = imrotate(image, -rotation, 'bilinear', 'crop');  
    [responses, temp_max_scales] = ...  
        multiscale_correlation(rotated, template, scales);  
    responses = imrotate(responses, rotation, 'nearest', 'crop');  
    temp_max_scales = imrotate(temp_max_scales, rotation, ...  
                               'nearest', 'crop');  
    higher_maxes = (responses > max_responses);  
    max_responses(higher_maxes) = responses(higher_maxes);  
    max_scales(higher_maxes) = temp_max_scales(higher_maxes);  
    max_rotations(higher_maxes) = rotation;  
end
```

Review: Correlation-based Detection

```
function result = find_template(image, template, scales, ...
                                rotations, result_number)

% function result = find_template(image, template, scales)
%
% returns the bounding boxes of the best matches for the template in
% the image, after searching all specified scales and rotations.
% result_number specifies the number of results (bounding boxes).

[max_responses, max_scales, max_rotations] = ...
    template_search(image, template, scales, ...
                    rotations, result_number);

result = detection_boxes(image, template, max_responses, ...
                        max_scales, result_number);
```

Review: Correlation-based Detection

```
function [result, boxes] = template_detector_demo(image, template, ...
                                                scales, rotations, result_number)

% function [result, boxes] =
%         template_detector_demo(image, template, ...
%         scales, rotations, result_number)
%
% returns an image that is a copy of the input image, with
% the bounding boxes drawn for each of the best matches for
% the template in the image, after searching all specified
% scales and rotations.

boxes = find_template(image, template, scales, ...
                    rotations, result_number);
result = image;

for number = 1:result_number
    result = draw_rectangle1(result, boxes(number, 1), ...
                            boxes(number, 2), ...
                            boxes(number, 3), boxes(number, 4));
end
```

Including Information From PCA

- One approach: call `pca_score` on every possible window.
 - Drawback: slow:
- How can we combine a slow but more accurate approach and a fast but less accurate approach?

Filter-and-refine Approach

- How can we combine a slow but more accurate approach and a fast but less accurate approach?
- **Filter step:** use the fast approach to get a shortlist of candidates.
- **Refine step:** use the slow approach to rerank the shortlist and choose the final results.

Filter-and-refine Face Detection

- **Filter step:** use correlation to identify a number of candidates.
 - number of candidates is 50 in the code.
- **Refine step:** use `pca_score` to evaluate those candidates.

```

function result = find_faces(image, scales, result_number, ...
                             eigenface_number)

load face_filter;
load final_eigens;

% start filter step
preliminary_number = 50;
preliminary_results = find_template(image, face_filter, ...
                                    scales, 0, preliminary_number);

% start refine step
for number = 1:preliminary_number
    top = preliminary_results(number, 1);
    bottom = preliminary_results(number, 2);
    left = preliminary_results(number, 3);
    right = preliminary_results(number, 4);
    window = image(top:bottom, left:right);
    preliminary_results(number, 6) = ...
        pca_score(window, mean_face, eigenvectors, eigenface_number);
end

[values, indices] = sort(preliminary_results(:, 6), 'ascend');
top_indices = indices(1: result_number);
result = preliminary_results(top_indices, :);

```

Recap on PCA for Face Detection

- What model for faces are we using?

Recap on PCA for Face Detection

- What model for faces are we using?
 - Faces are reconstructed well using the top K eigenvectors.
 - NOTE: each K defines a different model.
- What is a perfect face under this model?

Recap on PCA for Face Detection

- What model for faces are we using?
 - Faces are reconstructed well using the top K eigenvectors.
 - NOTE: each K defines a different model.
- What is a perfect face under this model?
 - A face spanned using the top K eigenvectors.
 - A face F such that:
 - `pca_backprojection(pca_projection(F)) = F.`
 - A face on which `pca_score` returns 0.

A Generalized View of Classifiers

- We have studied two face detection methods:
 - Normalized correlation.
 - PCA.
- Each approach exhaustively evaluates all image subwindows.
- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.

Features and Classifiers

- Our goal, in the next slides, is to get a better understanding of:
 - What is a feature?
 - What is a classifier?

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is a feature here?

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is a feature here?
- Two possible answers:
 - Each pixel value is a feature.
 - The feature is the result of normalized correlation with the template.

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is the score of each subwindow?

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is the score of each subwindow?
 - The result of normalized correlation with the template.
 - Arguably, the score is the feature itself.

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- How does the decision depend on the score?

Normalized Correlation Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- How does the decision depend on the score?
 - In `find_template`, faces are the top N scores.
 - N is an argument to the `find_template` function.
 - An alternative, is to check if $\text{score} > \text{threshold}$.
 - Then we must choose a threshold instead of N.

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is a feature here?

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is a feature here?
 - The result of the `pca_projection` function.
 - In other words, the K numbers that describe the projection of the subwindow on the space defined by the top K eigenfaces.

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is the score of each subwindow?

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- What is the score of each subwindow?
 - The result of the `pca_score` function.
 - The sum of squared differences between:
 - the original subwindow W , and
 - `pca_backprojection(pca_projection(W))`.

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- How does the decision depend on the score?

PCA-based Detection Analysis

- Each subwindow is evaluated in three steps:
 - First step: extract *features*.
 - Feature: a piece of information extracted from a pattern.
 - Compute a score based on the features.
 - Make a decision based on the score.
- How does the decision depend on the score?
 - In `find_faces`, faces are the top N scores.
 - N is an argument to the `find_faces` function.
 - An alternative, is to check if $\text{score} > \text{threshold}$.
 - Then we must choose a threshold instead of N.

Defining a Classifier

- Choose features.
- Choose a scoring function.
- Choose a decision process.
 - The last part is usually straightforward.
 - We pick the top N scores, or we apply a threshold.
- Therefore, the two crucial components are:
 - Choosing features.
 - Choosing a scoring function.

What is a Feature?

- Any information extracted from an image.
- The number of possible features we can define is enormous.
- Any function F we can define that takes in an image as an argument and produces one or more numbers as output defines a feature.
 - Correlation with a template defines a feature.
 - Projecting to PCA space defines features.
 - More examples?

What is a Feature?

- Any information extracted from an image.
- The number of possible features we can define is enormous.
- Any function F we can define that takes in an image as an argument and produces one or more numbers as output defines a feature.
 - Correlation with a template defines a feature.
 - Projecting to PCA space defines features.
 - Average intensity.
 - Std of values in window.